

## AY19/20 Semester 2 : DSA4212 Python Tutorials

### matplotlib

---

```
import matplotlib as plt

plt.plot(loss_history, "-.") # dotted-dash line
plt.plot(loss_history, "-x") # line with each point
                              # marked as 'x'
plt.plot(loss_history, "-s") # line with each point
                              # marked with a filled square
plt.plot(loss_history, "--x") # dotted lines with
                              # each point marked as 'x'

plt.plot(x_list, y_list, "r-", linewidth = 3) #
        # thick red line
plt.plot(x_list, z_MLE, "g-", linewidth = 3,
        label="Reconstruction") # labelled thick green
        # line
plt.plot(x_list, y_list, "bo", alpha = 0.5) # big
        # blue filled dots, with some transparency
        # black: 'k'

# Axis scales
plt.yscale("log")

plt.grid(True)

# Axis labels
plt.xlabel("iteration")
plt.ylabel("loss")

# Plot 2 plots side by side
plt.figure(figsize = (10, 5))

plt.subplot(1, 2, 1)
```

```
plt.plot(hloss_history)
plt.xlabel("iter")
plt.ylabel("horizontal loss")
# Labels involving math symbols
plt.xlabel(r"$\mu$")
plt.ylabel(r"$\sigma$")
plt.grid(True)

plt.subplot(1, 2, 2)
plt.plot(vloss_history)
plt.xlabel("iter")
plt.ylabel("vertical loss")
plt.grid(True)

# Plotting level lines
x = np.linspace(0, 2, num = 100)
y = np.linspace(0, 2.2, num = 100)
X, Y = np.meshgrid(x, y)
Z = func((X, Y))
plt.contour(X, Y, Z, colors='black')

# Scatterplot
plt.scatter(x_data, y_data)
plt.scatter(samples_gaussian[:,0],
            samples_gaussian[:,1], alpha=0.2, marker="+") #
            # many points
plt.scatter(mu_MLE, sigma_MLE, c="red", marker="^",
            s=100, label="MLE") # one point (filled red
            # triangle)
plt.scatter(mu_true, sigma_true, c="purple",
            marker="s", s=100, label="True") # one point
            # (filled purple square)
```

```
# Plot multiple lines and points, label and with
        # legend
plt.plot(x_list, z_true, "r--", label="Truth")
plt.plot(x_list, z_MLE, "g-", linewidth = 3,
        label="Reconstruction")
plt.plot(x_list, y_list, "bo", alpha=0.3,
        label="data")
plt.legend()
```

---

### Reading in Data

---

```
# Read in csv file
path = "../datasets/"
file = "abalone.data"
pd.read_csv(os.path.join(path, file), sep = ",",
            header = None) # "../datasets/abalone.data"

# Read in image
M = imageio.imread("nus_logo.png").astype(float)
M = onp.mean(M, axis=2) / 255. # transform to
        # greyscale
d = 100
M = resize(M, (d,d)) # resize
plt.imshow(M, cmap="gray") # display
```

---

## Others

---

```
# Type of object
type(obj)

# Selecting rows and columns in pandas
attribute.iloc[:20000, [0,-1]] # first 20000 rows,
                                first and last cols

# Similar to R which() for indices of True in
  boolean arrays
np.where(arr == 1)[0]

# Create a sequence of numbers
X = np.linspace(1.4, 1.9, num = 10) # R: seq(from,
  to, length)
Y = np.arange(2, 2.15, 0.02) # R: seq(from, to, by)

# Combine arrays pairwise
Z = np.array((X, Y)).T # X,Y same length

# Apply function over values in array
np.array(list(map(func, Z))) # over rows

# Sample random integers from a given range
randint(low, high, size) # inclusive end points
# Sample of random real numbers
np.random.rand(n)
# Sample values from given array with specified
  probabilities
np.random.choice(arr, 5, p = [0.5, 0.1, 0.1, 0.3],
  replace = False)

# Combine two arrays into one
np.array(arr1, arr2)

# Mean of values in array of arrays
np.mean(arr, axis = 0) # along cols
np.mean(arr, axis = 1) # along rows

# Mapping function over array(s)
def func(a, b, c):
```

```
    return a + b + c
```

```
A = np.arange(10) # [0, 1, 2, ..., 9]
B = np.arange(10)
C = np.arange(10)

mapped_func = jax.vmap(func, in_axes=(0, 0, 0)) #
  OR: in_axes(None, 0, 0)
mapped_func(A, B, C) # [ 0, 3, 6, 9, 12, 15, 18,
  21, 24, 27]

# Convert pandas Series object to NumPy array
obj.to_numpy(dtype="float64")
# Convert type of data in array of values
onp.array(x_data).astype(float)

# Concat NumPy arrays
A = np.array((1, 2, 3))
B = np.array((4, 5, 6))
C = np.array((7, 8, 9))

np.vstack((A, B, C)) # [ [1, 2, 3],
  #   [4, 5, 6],
  #   [7, 8, 9] ]
np.hstack((A, B, C)) # [1, 2, 3, 4, 5, 6, 7, 8, 9]

# Change map obj to array
np.array(list(map_obj)) # takes a while for large
  dim

# Set seed
import random
random.seed(n)

# Cumulative summation of an array of values
np.cumsum(arr)

# Mathematical functions
np.sin(theta)
np.cos(theta)

# Flatten a matrix (to a long vector)
x.flatten()
```

```
x.ravel() # faster than flatten()
```

```
# Distribution
onp.random.poisson(theta)
```

---

## Mathematics

---

```
# Factorial, n!
math.factorial(n)
```

---

# Logistic regression

**Loss function:**  $\frac{1}{N} \sum_{i=1}^N \log(1 + \exp(-y_i \langle \beta, x_i \rangle))$

**Logistic probability:**  $P(y = 1|x) = \frac{1}{1 + \exp(-\langle \beta, x \rangle)}$

Discussed in Tutorial 3 Exercise 4 - *fit logistic regression (i.e. find  $\beta$ ) with basic gradient descent by minimising loss function.*

Download the fashion dataset and load them into a jupyter notebook. The labels +1 and -1 correspond to whether a piece of clothing has long or short sleeve. Implement a logistic regression on the training dataset and train it with a basic gradient descent algorithm. Evaluate the accuracy of your model on the test dataset. **Remark:** It may be useful to learn how to use the vmap function from JAX.

```
# load libraries
import jax.numpy as np
import numpy as onp
import os as os
import pylab as plt
import h5py # to read hdf5 files

# load datasets
path = "../datasets/"

filename = "fashion_img_training.h5" # 10 000
    samples, each of size 28x28 (train set)
fashion_img_training =
    h5py.File(os.path.join(path,
        filename), 'r')['training_images'][:]

filename = "fashion_label_training.h5" # array of
    1, -1s
fashion_label_training =
    h5py.File(os.path.join(path,
        filename), 'r')['training_labels'][:]

filename = "fashion_img_val.h5" # 2 000 samples
    (validation set)
```

```
fashion_img_val = h5py.File(os.path.join(path,
    filename), 'r')['val_images'][:]

filename = "fashion_label_val.h5"
fashion_label_val = h5py.File(os.path.join(path,
    filename), 'r')['val_labels'][:]

# plot a sample of images to view
plt.figure(figsize=(15,1))
for k in range(15):
    plt.subplot(1,15,k+1)
    plt.imshow(fashion_img_training[k],
        cmap="gray")
    plt.title(str(fashion_label_training[k]))
    plt.axis("off")
```



Figure 1: Long sleeve (1.0), short sleeve (-1.0)

```
# define probability and loss functions
def proba_single(beta, x):
    """ beta is a vector of dimension 784, and x
        as well """
    proba = 1. / (1. + np.exp(-np.dot(beta, x)))
    return proba

# returns probabilities of 1.0 on a whole dataset
prediction_data = jax.vmap(proba_single,
    in_axes=(None,0))

def loss_single(beta,x,y):
    """
    beta: vector of dimension 784
    x: vector of dimension 794
    y: a number that equals -1 or 1
    """
    return np.log(1. + np.exp(-y * np.dot(beta,
        x)))
```

```
# compute the loss on. whole dataset
loss_dataset = jax.vmap(loss_single,
    in_axes=(None,0,0))

def loss(beta, data, y):
    """ compute the mean of all the individual
        losses """
    list_of_all_losses = loss_dataset(beta, data,
        y)
    return np.mean(list_of_all_losses)

def accuracy(beta, data, y):
    pred = prediction_data(beta, data)
    threshold = 0.5
    prediction_binary = (pred > 0.5).astype(int)
    #equals 1 if pred > threshold and 0
    #otherwise
    prediction_sign = 2*(prediction_binary - 0.5)
    #equals 1 if pred > threshold and -1
    #otherwise
    return np.mean(prediction_sign == y)

# initialise values
beta_init = onp.random.normal(0, 0.1, size=784)
grad_loss = jax.grad(loss, argnums=0)
grad_loss = jax.jit(grad_loss) #compile it to make
    it faster
beta = onp.copy(beta_init)

n_iter = 500
learning_rate = 0.1
loss_history = []
acc_train_list = []
acc_val_list = []

# gradient descent to find optimal beta, min loss
for k in range(n_iter):
    gradient = grad_loss(beta, data_train, y_train)
    beta = beta - learning_rate * gradient # update
    current_loss = loss(beta, data_train, y_train)
    # value to min
    accuracy_train = accuracy(beta, data_train,
        y_train)
```

```

accuracy_val = accuracy(beta, data_val, y_val)

acc_train_list.append(accuracy_train)
acc_val_list.append(accuracy_val)
# prints e.g. "Loss:2.014 Accuracy: 50.000% /
50.850%"
if k % 50 == 0:
    print("Loss:{0:.3f} \t
    Accuracy: {1:.3f}% /
    {2:.3f}%".format(current_loss, \
    100*accuracy_train,\
    100*accuracy_val))
loss_history.append(current_loss)

```

```

# plot loss and accuracy of train and val sets
plt.figure(figsize=(10,5))
plt.subplot(1,2,1)
plt.plot(loss_history)
plt.grid(True)
plt.title("Loss")
plt.xlabel("Iteration")

plt.subplot(1,2,2)
plt.plot(acc_train_list)
plt.plot(acc_val_list)
plt.grid(True)
plt.title("Accuracy")
plt.xlabel("Iteration")

```




Figure 2: Loss and accuracy of basic GD

Highlight beta weights of pixels which are more than 0.2, notice that these are mostly focused on the sleeve area.

```

plt.figure(figsize=(15,1))
for k in range(15):
    plt.subplot(1,15,k+1)
    plt.imshow(fashion_img_training[k],
    cmap="gray")
    plt.imshow(beta.reshape((28,28)) > 0.2,
    alpha=0.7)
    plt.title(str(fashion_label_training[k]))
    plt.axis("off")

```




Figure 3: Beta weights with more than 0.2

It is good to try and interpret the model. Beta  $\beta$  can be randomly initialised with a small variance (e.g. 0.1) with mean 0. Logistic regression is convex optimisation, so its easy in this case even with high dimension.

## Regularisation

### → Ridge regression

Minimise the function

$$L(\beta) = \frac{1}{N} \sum_{i=1}^N \log[1 + e^{-y_i \langle x_i, \beta \rangle}] + \frac{1}{N} \|\beta\|^2$$

Discussed in Tutorial 5 - *minimise logistic function with ridge regularisation using stochastic gradient descent.*

### → LASSO

Minimise the function

$$L(\beta) = \sum_{i=1}^N \log[1 + e^{-y_i \langle x_i, \beta \rangle}] + \lambda \|\beta\|$$

## Gradient Descent

The optimisation problem and goal is to **minimise** or **maximise** a function.

## Naive gradient descent

$$x_{n+1} = x_n - \eta \nabla f(x_n)$$

Discussed in Tutorial 3 Exercise 3 (Rosenbrock function) - *minimise given function with gradient descent.*

Consider the function

$$F(x, y) = (x - 1)^2 + 10(y - x^2)^2$$

*Context:* The Rosenbrock function is **non-convex** and often used as test problem for optimisation algorithms to see if it works well. The general form of the function is

$$F(x, y) = (a - x)^2 + b(y - x^2)^2$$

with a global minimum of 0 at  $(a, a^2)$ . Usually parameters are set such that  $a = 1$  and  $b = 100$ . When  $a = 0$ , the function is symmetric and the minimum is at the origin.

1. Implement a basic gradient descent algorithm to minimise  $F$ .

```

# import packages
import jax.numpy as np
import jax
import matplotlib as plt

# define rosenbrock function to minimise
def func(x):
    """
    x = (x[0], x[1])
    """
    return (x[0]-1)**2 + 10*(x[1] - x[0]**2)**2

# initialise parameters

```

```

x_init = np.array([2., 2.]) # x0
grad_func = jax.grad(func)
n_iter = 100
learning_rate = 0.01

loss_history = [] # monitor changes in func(x)
                    values
x1_history = [] # monitor changes in x values
x2_history = []
x = x_init[:]

for k in range(n_iter):
    gradient = grad_func(x)
    x = x - learning_rate * gradient
    x1_history.append(x[0])
    x2_history.append(x[1])
    loss_history.append(func(x))

# plot function values
plt.plot(loss_history, "-.") # dotted-dash line
plt.yscale("log")
plt.grid(True)
plt.xlabel("iteration")
plt.ylabel("Loss")

```




Figure 4: Plot of  $x_n$  across iterations

- Find a way to visualise the trajectory of the algorithm, as well as the level lines of the function  $F$ .

```

x = np.linspace(0, 2, num = 100)
y = np.linspace(0, 2.2, num = 100)
X, Y = np.meshgrid(x, y)
Z = func((X, Y))

plt.plot(x1_history, x2_history) # trajectory of x
                                values
plt.contour(X, Y, Z, colors='black') # level lines
plt.plot(1, 1, "ro") # true minimum

```




Figure 5: Trajectory of algorithm

## Stochastic gradient descent

Refer to section *Momentum* for code example.

## Descent algorithm

$$x_{n+1} = x_n + \eta_n d_n$$

where descent direction  $d = -S\nabla F(x)$

s.t.  $\langle d, \nabla F(x) \rangle < 0$

for any PD matrix  $S$  (where  $S = I_d$  for naive GD)

### → Backtracking

Set  $\alpha \in (0, 1)$  (typically 0.5), find step size  $\eta_n > 0$  such that  
Backtracking cond:  $F(x_n + \eta d_n) < F(x_n) + \alpha \eta \langle \nabla F(x_n), d_n \rangle$

Algo: Start from initial guess  $\eta$ , keep dividing it by 2 until condition satisfied (use this  $\eta$  to do update).

A variant of Tutorial 3 Exercise 3 (Rosenbrock function)  
- minimise given function with naive gradient descent **plus backtracking**.

```

x_init = np.array([5., 5.])
grad_func = jax.jit(jax.grad(func))

init_learning_rate = 1.0 # too large may not
                        decrease fn in each step
n_iter = 100
alpha = 0.5 # does not change
x = x_init
loss_history = []
learning_rate_history = []

for k in range(n_iter):
    # compute a descent direction
    gradient = grad_func(x)
    descent = -gradient

```

```

# keep reducing the learning rate until
# backtracking condition is satisfied
# for each iteration, learning rate starts
# again from init value
learning_rate = init_learning_rate
while func(x + learning_rate * descent) >
    func(x) +
    alpha*learning_rate*np.dot(gradient,
    descent):
    learning_rate = learning_rate / 2.

# actually do the update
x = x + learning_rate * descent

# save the history of loss / learning-rates
loss_history.append(func(x))
learning_rate_history.append(learning_rate)

# plot
plt.figure(figsize=(10,5))

plt.subplot(1,2,1)
plt.plot(loss_history, "-s")
plt.yscale("log")
plt.title("Loss vs iteration")
plt.grid(True)

plt.subplot(1,2,2)
plt.plot(learning_rate_history)
plt.title("Learning Rate vs iterationr")
plt.grid(True)

```

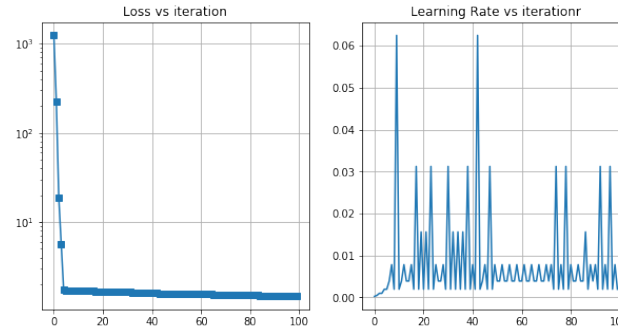


Figure 6: Loss and learning rate (Version 1)

**OR: a slightly different implementation (con't from previous iteration)**

```

for k in range(n_iter):
    # compute a descent direction
    descent = -grad_func(x)
    gradient = grad_func(x)

    # keep reducing the learning rate until
    # backtracking condition is satisfied
    # *DIFFERENT*
    # learning rate continues from previous
    # iteration
    learning_rate = learning_rate*3 # prevent
    # value from becoming too small
    while func(x + learning_rate * descent) >
        func(x) +
        alpha*learning_rate*np.dot(gradient,
        descent):
        learning_rate = learning_rate / 2.

    # actually do the update
    x = x + learning_rate * descent

    loss_history.append(func(x))
    learning_rate_history.append(learning_rate)

```




Figure 7: Loss and learning rate (Version 2)

**OR: Automatic backtracking in logistic regression model**

Discussed in Tutorial 4 Exercise 3 (Backtracking) - *logistic regression with backtracking*.

Consider the logistic regression model used in Tutorial 3 for differentiating between two types of clothes. Re-implement this algorithm using the automatic backtracking approach for choosing the gradient step-size.

```

n_iter = 100
learning_rate = 10.
loss_history = []
acc_train_list = []
acc_val_list = []
alpha = 0.5

for k in range(n_iter):
    # descent direction
    gradient = grad_loss(beta, data_train, y_train)
    descent_direction = -gradient

    # adaptive learning rate
    learning_rate = learning_rate * 3
    current_loss = loss(beta, data_train, y_train)
    while loss(beta + learning_rate *
        descent_direction, data_train, y_train) >
        current_loss +
        alpha*learning_rate*np.dot(descent_direction,

```

```

gradient):
learning_rate = learning_rate / 2.

#actually do the update
beta = beta + learning_rate * descent_direction

current_loss = loss(beta, data_train, y_train)
accuracy_train = accuracy(beta, data_train,
y_train)
accuracy_val = accuracy(beta, data_val, y_val)

acc_train_list.append(accuracy_train)
acc_val_list.append(accuracy_val)
if k % 10 == 0:
    print("Loss:{0:.3f} \t Accuracy: {
1:.3f}% / {2:.3f}%".format(current_loss, \
100*accuracy_train,\
100*accuracy_val))
loss_history.append(current_loss)

```

## → Exact line search

Find step-size  $\eta_n$  such that

$$F(x_n + \eta_n d_n) = \min_{\eta > 0} F(x_n + \eta d_n)$$

## Learning rate

### → Fixed

In each update of the parameter (e.g.  $x, \beta$ ), learning rate is fixed and remains the same.

### → Adaptive

1. Backtracking (find suitable step-size with backtracking condition)

## Momentum

$$x_{n+1} = x_n - \eta \nabla F(x_n) + \beta(x_n - x_{n-1})$$

for a fixed  $\beta > 0$  (usually 0.9) where we go in roughly the same direction as before.

Discussed in Tutorial 5 (Stochastic GD with momentum)  
- *fit regularised logistic regression of fashion dataset with stochastic GD with momentum*

2. Compute the Hessian for an arbitrary vector  $\beta$  and check that it is positive definite.

```

# random initialization
n_training_data, D = data_train.shape
beta_init = onp.random.normal(0, scale =
1./onp.sqrt(D), size=D)

# visualise the Hessian
H = hess_loss(beta_init, data_train, y_train)
eigenvalues, _ = onp.linalg.eigh(H) # all positive
plt.imshow(H)

```




Figure 8: Hessian heatmap

Use a moving average of stochastic gradient estimates.

For a moving average of  $m \approx 0.9$ , compute the approximated gradient to use in each update as

$$g_k = m g_{k-1} + (1 - m) \hat{\nabla} L(\beta)$$

This is SGD of logistic regression with momentum and adaptive learning rate.

```

# initialise values
n_epoch = 50
batch_size=100 # size of the mini-batch (SGD)
beta = onp.copy(beta_init)
mov_avg_param = 0.95

learning_rate = 0.1
loss_train_history = []
loss_val_history = []

acc_train_history = [] # computed at the end of
each epoch
acc_val_history = []

# initialize gradient
gradient = grad_loss(beta, data_train, y_train)

start = time.time()
for epoch in range(n_epoch):
    # divide learning rate by 2 every 10 epoch
    if epoch % 10 == 0 and epoch >= 1:
        learning_rate = learning_rate/2.

    for k in range(n_training_data // batch_size):
        # number of batch per epoch

        # 1st option
        #=====
        # select at random "batch_size" random data
        point
        # index_ =
        onp.random.choice(train_data_size,
batch_size, replace=False)

        # 2nd option

```

```

#=====
# consider the batches sequentially
index_ = onp.arange(k*batch_size,
                    (k+1)*batch_size) % n_training_data
gradient_local = grad_loss(beta,
                           data_train[index_], y_train[index_])

# moving average of the gradient ***
gradient = gradient * mov_avg_param + (1 -
                                       mov_avg_param) * gradient_local

# gradient descent update
beta = beta - learning_rate * gradient

# before moving on to next epoch
accuracy_train = accuracy(beta, data_train,
                           y_train)
accuracy_val = accuracy(beta, data_val, y_val)
acc_train_list.append(accuracy_train)
acc_val_list.append(accuracy_val)

loss_train = loss(beta, data_train, y_train)
loss_val = loss(beta, data_val, y_val)
loss_train_history.append(loss_train)
loss_val_history.append(loss_val)

timing = time.time() - start
print("epoch:{0:.0f} \t time:{1:.0f} \t
      Loss(train):{2:.3f} \t Loss(val):{3:.3f}"
      .format(epoch, timing, loss_train, loss_val))

# plotting
plt.plot(loss_train_history, label="loss(train)")
plt.plot(loss_val_history, label="loss(val)")
plt.legend()
plt.grid(True)
plt.xlabel("epoch")
plt.title("Stochastic Gradient Descent")

```

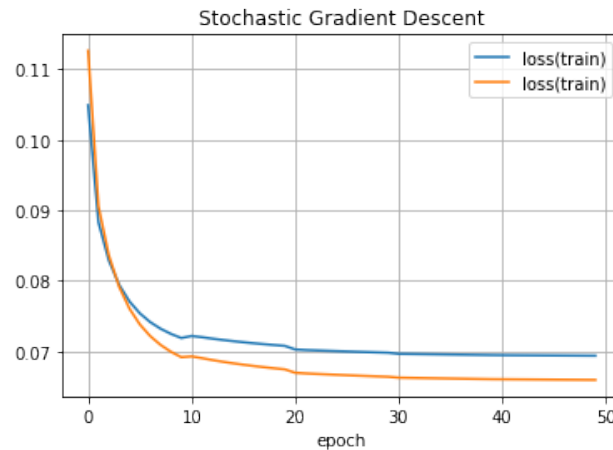


Figure 9: Higher curve (train), lower curve (val)

## Pre-conditioning

$$x_{k+1} = x_k - \eta_k P_k \nabla F(x_k)$$

$$\text{OR: } x_{k+1} = x_k - P_k d_k$$

$$\text{Descent direction: } d_n = P \nabla F(x_n)$$

Discussed in Tutorial 4 Exercise 4 (Pre-conditioning) - *minimise given function with preconditioned gradient descent with backtracking*

Consider a positive definite symmetric matrix  $S \in \mathbb{R}^d$ ,  $d = 100$ , and the quadratic function

$$F(x) = (1/2) \langle x, Sx \rangle$$

1. What is the minimum of the function  $F$ ? *Ans: 0, when  $x$  is 0 and given  $S > 0$*
2. In this exercise, we will take  $S = C^{10}$  where  $C$  is the empirical covariance matrix of  $N = 100$  random Gaussian vectors. Implement this in Python.

---

```

# Find cov mat C
N = 300
d = 100
gaussian_samples = onp.random.normal(0, 1, size =
                                     (N, d))
empirical_cov = onp.cov(gaussian_samples.T) # cols
                                     are obs
S = onp.copy(empirical_cov)
# for C^10
for k in range(9):
    S = S @ empirical_cov

```

---

3. Implement gradient descent with backtracking to minimise  $F$ . *Ans: refer to backtracking*
  4. Implement the Preconditioned Gradient Descent (PGD) with backtracking.
- 

```

def func(x):
    return 0.5 * np.dot(x, np.dot(S, x))
grad_func = jax.jit(jax.grad(func))

# PGD
def run_gradient_descent(P, n_iter,
                        learning_rate_init, x_init):
    """
    P: preconditioning matrix
    n_iter: number of iteration
    learning_rate_init: initial learning rate
    x_init: initial guess
    """

    loss_history = []
    alpha = 0.5
    learning_rate = learning_rate_init

    x = onp.copy(x_init)
    for k in range(n_iter):
        # descent direction
        gradient = grad_func(x)
        descent_direction = -P @ gradient # NOTE:
        not gradient

```



```

# adaptive learning rate (backtracking)
learning_rate = learning_rate * 3
current_loss = func(x)
while func(x + learning_rate *
    descent_direction) > current_loss +
    alpha*learning_rate*np.dot(descent_direction,
    gradient):
    learning_rate = learning_rate / 2.

# actually do the update
x = x + learning_rate * descent_direction
current_loss = func(x)
loss_history.append(current_loss)

return loss_history

```

5. Implement PGD with pre-conditioning matrices  $P = S$ ,  $P = I_d$  and  $P = S^{-1}$ . Which one converges faster?

```

n_iter = 100
learning_rate_init = 10.
x_init = onp.random.normal(0,1,size=d)

# 1) P = S
P = onp.copy(S)
loss_history_P = run_gradient_descent(P, n_iter,
    learning_rate_init, x_init)

# 2) P = Id (basic GD)
P = onp.eye(d) # identity matrix
loss_history_identity = run_gradient_descent(P,
    n_iter, learning_rate_init, x_init)

# 3) P = inverse S
eps = 0.001
P = onp.linalg.inv(S + eps*onp.eye(d)) # make sure
    matrix invertible, add to eigenvalues, else
    difficult to converge
loss_history_P_inv = run_gradient_descent(P,
    n_iter, learning_rate_init, x_init)

```

```

# plot
plt.plot(loss_history_identity, label="naive")
plt.plot(loss_history_P, label="C")
plt.plot(loss_history_P_inv, label="C_inv")
plt.legend()

plt.grid(True)
plt.yscale("log")

plt.title("Comparison Preconditionning Matrix")
plt.ylabel("log loss")
plt.xlabel("Iteration")

```

Choice of pre-conditioning matrix does matter, it can speed up the algorithm dramatically.

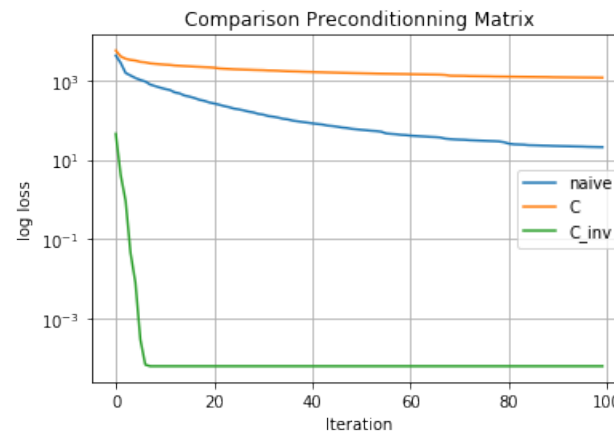


Figure 10: Comparison of pre-conditioning matrices

## Newton method

$$x_{k+1} = x_k - \eta_k H_k^{-1} \nabla F(x_k)$$

An alternative to computing inverse of Hessian is to solve the equation to find the descent direction .

$$H_k d_k = -\nabla F(x_k)$$

Discussed in Tutorial 5 - *fit regularised logistic regression of fashion dataset with Newton- algorithm.*

4. Implement a Newton-algorithm (with backtracking) for fitting the model.

```

# functions
...
grad_loss = jax.jit(jax.grad(loss))
hess_loss = jax.jit(jax.hessian(loss))
...

# random initialization
n_training_data, D = data_train.shape
beta_init = onp.random.normal(0, scale =
    1./onp.sqrt(D), size=D)

# initialise values
n_iter = 20
learning_rate = 0.9
loss_history = []
beta = onp.copy(beta_init)

for k in range(n_iter):
    # compute the gradient
    gradient = grad_loss(beta, data_train, y_train)

    # compute the Hessian
    H = hess_loss(beta, data_train, y_train)

    # no need to compute the inverse of H
    # better is to solve the system ***
    # H*d = -gradient
    descent_direction = onp.linalg.solve(H,
        -gradient)

```

```

# update estimate of the minimum
beta = beta + learning_rate * descent_direction

# save the loss for later plotting
current_loss = loss(beta, data_train, y_train)
loss_history.append(current_loss)

plt.plot(loss_history, "r-s")
plt.title("Newton Descent")
plt.xlabel("Iteration")
plt.ylabel("Loss")
plt.grid(True)

```

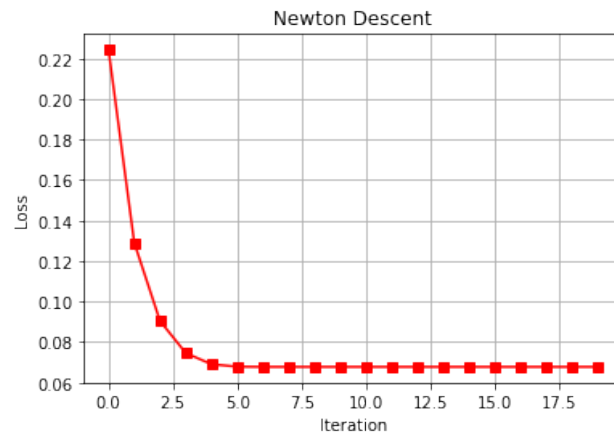


Figure 11: Trajectory of loss using Newton Descent across iterations

## Quasi-Newton methods

Approximate Hessian  $\mathbf{H}$  with  $B_k \approx \text{Hess}(x_k)$  .  
by solving for descent direction  $d_k$

$$B_k d_k = -\nabla F(x_k)$$

Important properties of approximated Hessian  $B_k$ :

1. Quasi-Newton methods can be applied for **any function** .
2.  $B_k$  guaranteed to be **PD** (even if function not convex).
3. **Finding** the approximated **inverse Hessian**  $B_k^{-1}$  is **fast** (in  $O(p^2)$  time).

For help in Python, look at `scipy.optimize.minimize?`

1. 'BFGS'
2. 'L-BFGS-B' : low memory item 'Nelder-Mead' : when you don't know how to compute the gradient

## Gauss-Newton methods

Used on functions which can be expressed as

$$F(x) = \ell(G(x))$$

where

$G : \mathbb{R}^P \rightarrow \mathbb{R}$  is a complicated function where only gradient  $G'$  can be computed and

$\ell : \mathbb{R} \rightarrow \mathbb{R}$  is a simple scalar function where first and second derivatives  $\ell', \ell''$  can be computed.

Gauss-Newton approximation to Hessian: .

$$\hat{H}(x) \approx \ell''(G(x)) \nabla G(x) \nabla G^T(x)$$

Differences of Gauss-Newton methods with Quasi-Newton methods:

1. Can only work on **composite functions** (i.e. expressed in the form  $\ell(G(x))$ ).
2. **Faster** (I think) at approximating Hessian of function  $F(x)$ .

→ **LBFGS**

**LBFGS** is short for Low-memory Broyden-Fletcher-Goldfarb-Shanno algorithm developed in 1973. **Note:** Using L-BFGS in Python, you may need to transform back JAX gradients into usual numpy vectors before feeding it to the L-BFGS function. **Also, looks like L-BFGS can only optimise a vector so flatten your images (arr of arr) and reshape it back in the loss function.**

Discussed in Tutorial 5 Example - *how to use LBFGS with scipy*.

```

# define function to minimise
def func(x):
    """ F: R^p --> R """
    return np.mean( np.sin(x)**2 )

# jax gradient

```

```

grad_func = jax.jit(jax.grad(func))

# create a wrapper, ensure the gradient output is
# standard numpy gradient
# should be faster than normal gradient descent
def gradient_wrapper(x):
    """ compute the gradient of func at x and make
    sure that the output is a numpy array"""
    return onp.array(grad_func(x))

# initialise values
dim = 100
x_init = onp.random.normal(0, 1, size = dim)

# optimisation
loss_history = [] # save the loss trajectory
time_history = [] # save the compute time

def save_traj(beta):
    """ a function that saves a few statistics for
    later analysis"""
    loss_history.append(func(beta))
    time_history.append(time.time() - start)

start = time.time()
traj = scipy.optimize.minimize(
    fun = func, #function to minimize
    x0 = onp.array(x_init), #initial guess
    method = 'L-BFGS-B',
    jac = gradient_wrapper, # function that computes
    the gradient
    callback = save_traj, # function used to save
    some results for later plotting
    options = {"maxiter": 100}, # maximum number of
    iterations
    #tol=10**-8
)
traj.x # IMPORTANT (optimal value found)

# plot results
plt.plot(time_history, loss_history, "r-s")
plt.xlabel("time (sec)")
plt.ylabel("loss")

```

```

plt.title("L-BFGS")
plt.grid(True)

```

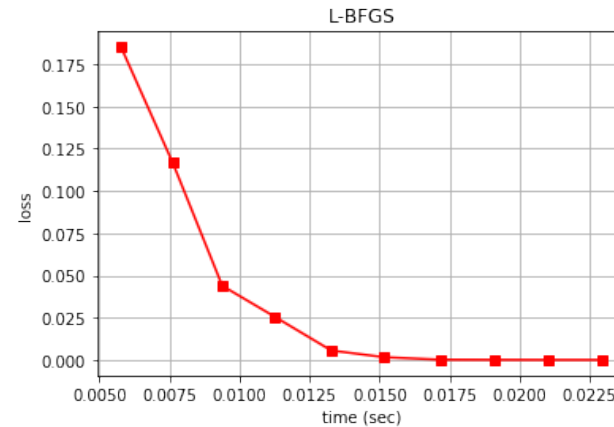


Figure 12: Trajectory of loss using L-BFGS across iterations (in terms of time)

**OR:** Discussed in Tutorial 5 - *fit regularised logistic regression of fashion dataset with LBFGS*.

5. Use the scipy implementation of L-BFGS for fitting the model.

```

# define function and gradient
def loss_wrapper(beta):
    """ only one argument """
    return loss(beta, data_train, y_train)

grad_loss_wrapper = jax.jit(jax.grad(loss_wrapper))

def gradient_wrapper(beta):
    """output is a numpy array"""
    return onp.array(grad_loss_wrapper(beta))

# optimise
loss_train_history = [] # save the loss trajectory

```

```

loss_val_history = []
time_history = [] # save the compute time

def save_traj(beta):
    """ a function that saves a few statistics for
    later analysis"""
    loss_train = loss(beta, data_train, y_train)
    loss_val = loss(beta, data_val, y_val)
    timing = time.time() - start
    print("Time:{0:3f} \t Loss(train):{1:.3f} \t
    Loss(val):{2:.3f}".format(timing, loss_train,
    loss_val))
    time_history.append(timing)
    loss_train_history.append(loss_train)
    loss_val_history.append(loss_val)

# run LBFGS
start = time.time()
traj = scipy.optimize.minimize(
    fun = loss_wrapper, # function to minimize
    x0 = onp.array(beta_init), # initial guess
    method = 'L-BFGS-B', # use L-BFGS
    jac = gradient_wrapper, # function that computes
    the gradient
    callback = save_traj, # function used to save
    some results for later plotting
    options={"maxiter":100}) # maximum number of
    iterations

```

Time:1.399521	Loss(train):0.200	Loss(val):0.218
Time:1.593403	Loss(train):0.183	Loss(val):0.196
Time:1.732385	Loss(train):0.175	Loss(val):0.187
Time:1.853453	Loss(train):0.138	Loss(val):0.148
Time:1.977781	Loss(train):0.104	Loss(val):0.109
Time:2.114910	Loss(train):0.092	Loss(val):0.095
Time:2.250759	Loss(train):0.082	Loss(val):0.084
Time:2.539433	Loss(train):0.078	Loss(val):0.077
Time:2.816728	Loss(train):0.075	Loss(val):0.073
Time:3.116674	Loss(train):0.073	Loss(val):0.069
Time:3.261371	Loss(train):0.071	Loss(val):0.066
Time:3.396481	Loss(train):0.069	Loss(val):0.065
Time:3.620170	Loss(train):0.069	Loss(val):0.065
Time:3.854951	Loss(train):0.069	Loss(val):0.065
Time:3.994454	Loss(train):0.068	Loss(val):0.064
Time:4.125747	Loss(train):0.068	Loss(val):0.064
Time:4.253343	Loss(train):0.068	Loss(val):0.064
Time:4.481123	Loss(train):0.068	Loss(val):0.064
Time:4.721010	Loss(train):0.068	Loss(val):0.064
Time:5.045678	Loss(train):0.068	Loss(val):0.064
Time:5.382473	Loss(train):0.068	Loss(val):0.064
Time:5.712736	Loss(train):0.068	Loss(val):0.064
Time:5.855145	Loss(train):0.068	Loss(val):0.064
Time:6.054924	Loss(train):0.068	Loss(val):0.064
Time:6.182820	Loss(train):0.068	Loss(val):0.064
Time:6.450254	Loss(train):0.068	Loss(val):0.064
Time:6.578574	Loss(train):0.068	Loss(val):0.064
Time:6.705868	Loss(train):0.068	Loss(val):0.064
Time:6.842521	Loss(train):0.068	Loss(val):0.064
Time:7.045810	Loss(train):0.068	Loss(val):0.064

Figure 13: Loss using L-BFGS across time

## Projected subgradient descent (Constrained optimisation)

Where subgradient  $g_n \in \partial f(x_n)$ ,

$$x_{n+1} = x_n - \eta_n g_n$$

$$x_{n+1} = \Pi_{\omega}(x_n - \eta_n g_n) \quad (\text{projected})$$

For a **subgradient** of  $f$  at  $x \in \mathbb{R}^p$ ,  
for subgradient vector  $g$ , for any vector  $\epsilon \in \mathbb{R}^p$ ,

$$f(x + \epsilon) \geq f(x) + \langle g, \epsilon \rangle$$

For a **subdifferential** of  $f$  at  $x \in \mathbb{R}^p$ ,

$$\partial f(x) = \{g \in \mathbb{R}^p : g \text{ is a subgradient to } f \text{ at } x\}$$

Note:

1. There exists **many subgradients** to function  $f$  at  $x \in \mathbb{R}^p$ .
2. A **subdifferential** of  $f$  at  $x$  is the **set of all possible subgradients** to  $f$  at  $x \in \mathbb{R}^p$ .
3. For **convex function**, subdifferential at  $x$  is always a **non-empty** and **bounded** set.
4. For **convex and differential at  $x$**  function, subdifferential contains only a **single vector** (i.e. gradient of function at  $x$ ,  $\partial f(x) = \{\nabla f(x)\}$ ).
5. Subgradient descent even works for non-differentiable convex function.

### → L2 ball constraint

Discussed in an example after Tutorial 6.

Minimise the function

$$F(x_1, x_2) = (x_1 - 3)^2 + (x_2 - 3)^2$$

under the constraint that  $(x_1, x_2) \in \omega$ , where  $\omega$  is a ball of radius  $R = 2$ .

---

```
def func(x):
    """ basic function """
    return (x[0]-3.)**2 + (x[1]-3.)**2

# compute a gradient
grad_func = jax.grad(func)

# IMPORTANT (constraint was norm of vector to be
# less than or equal to 2)
def proj_ball(x, R):
    """ projection on the disk on radius R """
    x_norm = np.linalg.norm(x)
    if x_norm <= R:
        return x
    else:
        return R * (x / x_norm) # on circumference
        of ball, same direction with unit
        vector
```

---

Continue with projected gradient descent. Save values before (after update) and after projection.

---

```
n_iter = 100
eta = 0.1

# store trajectory
trajectory_x = []
trajectory_y = []

# initialise values
R = 2
x_init = np.array([3., -3.])
x = x_init

# projected gradient descent
for k in range(n_iter):
    # gradient update
    g = grad_func(x)
    x = x - eta*g

# save it for plotting later
trajectory_x.append(x[0])
```

```

trajectory_y.append(x[1])

# do projection
x = proj_ball(x, R)

# save it for plotting later
trajectory_x.append(x[0])
trajectory_y.append(x[1])

```

Plot trajectory of values.

```

# plot circle (constraint space)
theta = np.linspace(0, 10, 1000)
plt.plot(R*np.sin(theta), R*np.cos(theta), "b-")

# plot the trajectory
plt.plot(trajectory_x, trajectory_y, "r-")

# plot the point (3,3) which is the unconstrained
# minimum (optimal point if no constraint)
plt.plot([3], [3], "gs")
plt.grid(True)
plt.title("Projected Gradient Descent")

```

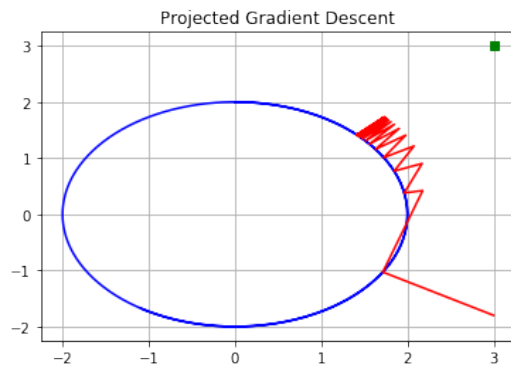


Figure 14: Trajectory of projected gradient descent with L2 ball as constrained subspace

## → Isotonic regression

Discussed in Tutorial 6 Question 2.

Consider the following synthetic dataset. Generate some random  $x_i$  and define

$$y_i = F(x_i) + (\text{noise})$$

for an unknown increasing function  $F$ . The goal of this exercise is to find an approximation of the (unknown) increasing function  $F$  from a set of samples  $\{x_i, y_i\}_{i=1}^N$ . In other words, this is again a regression problem, but this time with a constraint on the regression function.

```


n_data = 100
sd_noise = 0.4
x_data = onp.sort( 50 * onp.random.random(n_data)
) #generate and sort

def func(x):
    return onp.sqrt(x)

y_data = onp.sqrt(x_data) +
    onp.random.normal(loc=0, scale=sd_noise,
        size=n_data)
plt.scatter(x_data, y_data)
plt.plot(x_data, func(x_data), "r-")

```

First plot data with noise.



Our model again will be very simple. We will be looking

for  $z_1, \dots, z_N$  such that

$$z_i \approx y_i$$

and such that  $z_1 < z_2 < \dots < z_N$ . This means that  $z_i$  will represent the value of the unknown function  $F$  at location  $x_i$ ,

$$z_i = F(x_i)$$

In order to simplify things, we will instead be looking for non-negative coefficients  $\alpha_1, \dots, \alpha_N \geq 0$  and set

$$z_i = \alpha_1 + \dots + \alpha_i$$

Putting this all together, it means that we will be minimising the function

$$L(\alpha) \equiv \sum_{i=1}^N (y_i - [\alpha_1 + \dots + \alpha_i])^2$$

under the constraint that  $\alpha_1, \dots, \alpha_N \geq 0$ . Carry out this minimisation and plot the resultant function.


**Remark:** You might want to use the function `jax.numpy.cumsum`

```

# Introduce all increments
alpha = onp.random.rand(n_data)
f = np.cumsum(alpha)

plt.scatter(x_data, y_data)
plt.plot(x_data, func(x_data), "r-")
plt.plot(x_data, f, "b-")
plt.grid(True)

```



---

```
def loss(alpha):
    """
    alpha: set of increments
    """
    # compute function itself
    f = np.cumsum(alpha)
    # compute MSE
    return np.mean((y_data - f)**2)

grad_loss = jax.jit(jax.grad(loss))

def projection_positive(alpha):
    """ projection on the set of vectors with all
    coordinates positive """
    alpha_projected = onp.copy(alpha)
    # set to zero all coordinates that are negative
    alpha_projected[alpha_projected < 0] = 0
    return np.array(alpha_projected) # make sure
    array is JAX array

# Implement projected gradient descent algorithm
n_iter = 5000
loss_history = []
learning_rate = 0.001

for k in range(n_iter):
    gradient = grad_loss(alpha)

    # 1) gradient descent step
    alpha = alpha - learning_rate * gradient

    # 2) projection step
    alpha = projection_positive(alpha)

    # save loss
    loss_history.append(loss(alpha))

plt.plot(loss_history)
plt.yscale("log")
plt.grid(True)
```

---

With optimised values of  $\alpha$  (i.e. increments), plot out

fitted values with true values.

---

```
f_isotonic = np.cumsum(alpha)

plt.scatter(x_data, y_data, alpha = 0.5) # given
data points
plt.plot(x_data, func(x_data), "r-") # true
underlying function
plt.plot(x_data, f_isotonic, "b-", linewidth = 3)
```

---




Figure 15: Isotonic fitted regression (blue line) with true function (red line) and given data points (scatterplot)

The goal was to recover the blue line given the data points.

## Function Approximation

Discussed in Tutorial 6 Question 1 - given **mathematical formula** of function, approximate **parameter values of function**.

Consider the following synthetic dataset. Some  $x_i$  have been generated at random and defined


$$y_i = F(x_i) + (\text{noise})$$

The goal of this exercise is to find an approximation of the (unknown) function  $F$  from a set of samples  $\{x_i, y_i\}_{i=1}^N$ . In other words, this is a regression problem.

---

```
# Initialise data
n_data = 200
sd_noise = 0.7
x_data = onp.random.normal(loc=0, scale=5,
                             size=n_data)
# sine curve with normal noise
y_data = onp.sin(x_data) + onp.random.normal(loc=0,
                                                scale=sd_noise, size=n_data)
plt.scatter(x_data, y_data)
plt.grid(True)
```

---



The model used will be very simple. We would like to find

a function of the type

$$f(x; c) = \sum_{i=1}^N c_i \exp \left\{ -\frac{(x - x_i)^2}{2\ell^2} \right\}$$

for unknown coefficients  $c = (c_1, c_2, \dots, c_N)$ . For this exercise, you can take  $\ell = 0.2$  for example. In other words you will be minimising the function

$$f(x; c) = \sum_{i=1}^N (y_i - f(x_i; c))^2$$

Carry out this minimisation and plot the resulting function  $x \mapsto f(x; c^*)$ .

---

```
# Store unknown parameters in a vector of "c" of
# length n_data (i.e. c is like beta)
c = onp.random.normal(0, 1, size = n_data)

# Fix the lengthscale parameter "L"
LL = 1.

# Define function (NOT the function to minimise)
def func(x, c):
    return np.sum(c * np.exp( -(x - x_data)**2 /
        (2*LL) ))

# Mean squared loss function
def mse(c):
    # compute func(x_i) for each x_i and call the
    # output z
    z = np.array([func(x, c) for x in x_data])
    return np.mean((y_data - z)**2)

grad_mse = jax.grad(mse) # unfortunately, not
# "that" easy to JIT this function

# Wrapper to ensure output is the standard numpy
# gradient/array
def gradient_wrapper(c):
    return onp.array(grad_mse(c))
```

---

Minimise loss function with L-BGFS.

---

```
# Initialise guess
c_init = onp.random.normal(0, 1, size = n_data)

# Optimise with LBFGS
loss_history = [] # to save loss trajectory
def save_traj(c):
    curr_loss = mse(c)
    loss_history.append(curr_loss)
    print(curr_loss)

traj = scipy.optimize.minimize(
    fun = mse, # function to minimise
    x0 = onp.array(c_init) # initial guess
    method = 'L-BFGS-B',
    jac = gradient_wrapper, # fn to compute gradient
    callback = save_traj, # fn used to save some
    # results for later plotting
    options = {"maxiter": 50}
)
```

---

11.275136176213756	0.5066702437578349
5.231727920391103	0.5028709844938811
3.119553783660794	0.5000961848611705
1.3004176925787885	0.49621825335629294
1.1124141090873665	0.49271800365996
0.8074396758933589	0.48845729611575456
0.7467173830855871	0.4879480012956494
0.6976608086961855	0.48745575118553847
0.6591088803712636	0.48661345852458493
0.6195179054214998	0.4854206588271459
0.6136213007409879	0.4836848050129259
0.5986934942910623	0.48340008395813067
0.5935551620105362	0.4825185267729145
0.58533905800209	0.48199284308328044
0.5802464548690571	0.4811202482237931
0.5755808728257068	0.4804735509443418
0.573907308949833	0.4798102580541328
0.5712495769934665	0.479467058878419
0.5650542898654077	0.4793379834894317
0.5626467255391284	0.47918285361138613
0.5546652584398928	0.47848589118511603
0.5506785129414208	0.4780640510499224
0.5469578848553002	0.4776934109353638
0.5381090817162566	0.47738937165484074

---

```
# final estimate
c_lbfgs = traj["x"]

# draw estimated function
x_list = onp.linspace(-15, 15, 1000)
plt.plot(x_list, [func(x, c_lbfgs) for x in
    x_list], "-r", linewidth = 3)

# superimpose given data
plt.scatter(x_data, y_data)
plt.grid(True)
plt.title("Function Fitting")
```

---

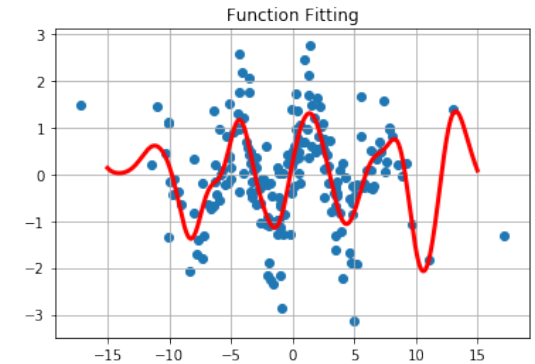


Figure 16: Approximated function with given data

## Support Vector Machine (SVM)

Discussed in Tutorial 6 Question 3.

In this question, we will consider the dataset `abalone.data` available at <https://archive.ics.uci.edu/ml/machine-learning-databases/abalone/>.

A SVM algorithm is a binary classifier similar to logistic regression and that can be described as follows: given training examples  $\{x_i, y_i\}_{i=1}^N$  with covariate  $x_i \in \mathbb{R}^p$  and binary response  $y_i \in \{-1, +1\}$ ; the SVM minimises with loss function

$$L(\beta) = \sum_{i=1}^N \max(0, 1 - y_i \langle \beta, x_i \rangle)$$

Once the SVM is fitted and the optimal parameter  $\beta_\star$  is found, the predicted value associated to a vector of covariates  $x$  simply equals the sign of the dot product  $\langle \beta_\star, x_i \rangle$ .

1. Download and load the dataset.
2. Use a SVM classifier to distinguish male from female abalone. To do so, first get rid of the training examples that correspond to infants (I).
3. Use 70% of the dataset to train a SVM classifier and estimate the accuracy on the remaining 30%. (**Remark:** It is normal if the accuracy is not great)

```
# Load dataset
path = "../datasets/"
file = "abalone.data"
data = pd.read_csv(os.path.join(path, file), sep =
    ",", header = None)

data
```

	0	1	2	3	4	5	6	7	8
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550	7
...	...	...	...	...	...	...	...	...	...
4172	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	11
4173	M	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	10
4174	M	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	9
4175	F	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	10
4176	M	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950	12

4177 rows x 9 columns

Figure 17: Quick look at abalone data

Form training and validation sets from given data.

```
sex = data.values[:, 0] # rename column name
x_data = data.values[(sex == "M") | (sex == "F"),
    1:] # select rows matching predicate, exclude
    first column (labels)
n, p = x_data.shape

# rescale covariates x_i s
for k in range(p):
    # select a column with x_data[:, k]
    x_data[:, k] = (x_data[:, k] - onp.mean(x_data[:,
        k])) onp.std(x_data[:, k])

x_data = onp.array(x_data).astype(float)
y_data = sex[(sex == "M") | (sex == "F")]
# Change labels to -1, +1
y_data[y_data == "M"] = 1.
y_data[y_data == "F"] = -1.
y_data = onp.array(y_data).astype(float)
```

Define loss functions, gradient and accuracy functions.

```
# Implement loss function for SVM
def loss_svm_single(beta, x, y):
```

```
    return np.maximum(0., 1. - y*np.dot(beta, x))

# Loss for entire dataset
loss_svm_all = jax.vmap(loss_svm_single, in_axes =
    (None, 0, 0))

# Global loss
def loss_svm(beta, x_batch, y_batch):
    return np.mean(loss_svm_all(beta, x_batch,
        y_batch))

# Gradient of loss
grad_svm = jax.jit(jax.grad(loss_svm))

def accuracy(beta, x_batch, y_batch):
    """
    proportion of correct predictions where prediction
    is 1 if the dot product is positive
    """
    return np.mean((x_batch @ beta > 0) == (y_batch >
        0))
```

Minimise loss function with gradient descent. Here, `x_data` is an array of arrays of size 2835 by 8.

```
# Random initialisation
beta_init = onp.random.normal(0, 1, size = p)

# Number of training data (70%)
n_train = int(len(x_data) * 0.7)

# Basic gradient descent (by default JAX computes a
    subgradient)
beta = onp.copy(beta_init)
learning_rate = 0.1
n_iter = 2000
accuracy_history = []
loss_history = []

for k in range(n_iter):
    gradient = grad_svm(beta, x_data[:n_train],
        y_data[:n_train]) # on train data
```



```

beta = beta - learning_rate * gradient

acc = accuracy(beta, x_data[n_train:],
                y_data[n_train:]) # on validation data
loss = loss_svm(beta, x_data[:n_train],
                y_data[:n_train])
loss_history.append(loss)
accuracy_history.append(acc)
if k % 200 == 0:
    print("iter:{} \t accuracy:{} \t
          loss:{}".format(k, acc, loss))

```

```

iter:0    accuracy:0.4606345475910693    loss:1.4060541813830736
iter:200  accuracy:0.5017626321974148    loss:0.9864938028325623
iter:400  accuracy:0.5381903642773208    loss:0.9638630691374537
iter:600  accuracy:0.5581668625146886    loss:0.9475591251239819
iter:800  accuracy:0.5569917743830788    loss:0.9365696482679869
iter:1000 accuracy:0.564042303172738    loss:0.9321896358774496
iter:1200 accuracy:0.5652173913043478    loss:0.9305409466388437
iter:1400 accuracy:0.5699177438307873    loss:0.929718042500022
iter:1600 accuracy:0.5710928319623971    loss:0.9292621797871887
iter:1800 accuracy:0.5710928319623971    loss:0.9289668670385685

```

Figure 18: Accuracy and loss values in each iteration of GD

```

plt.figure(figsize=(10,3))
plt.subplot(1,2,1)
plt.plot(loss_history)
plt.title("Loss")
plt.grid(True)

plt.subplot(1,2,2)
plt.plot(accuracy_history)
plt.title("Test Accuracy")
plt.grid(True)

```




Figure 19: Accuracy and loss plots for SVM algorithm

```

#display
plt.imshow(M, cmap="gray")

```




Figure 20: Read in image of NUS logo represented by a matrix of floats

Note that the larger the rank, the better the approximation while the lower the rank the lower the quality of the estimate since the image is being compressed.

### → Initialise values and setting up functions

Note that we have:  $U \in \mathbb{R}^{d,r}$  and  $V \in \mathbb{R}^{r,d}$

```

# Function and gradients
def loss(U,V):
    return np.mean( (M - U @ V)**2 )

# NOTE: how to do optimisation with 2 parameters
# differentiate gradient only with respect to
# parameter in question
grad_loss_u = jax.jit(jax.grad(loss, argnums = 0))
grad_loss_v = jax.jit(jax.grad(loss, argnums = 1))

# Init values
r = 20
U_init = onp.random.normal(0, 1, size = (d, r)) /
    onp.sqrt(d*r)

```

## Low Rank Approximation

Discussed in Tutorial 7 Question 3.

Load the image nus\_logo.png and transform it into greyscale and consider the associated matrix  $M \in \mathbb{R}^{d,d}$ . Let  $r \geq 1$  be an integer. Find two matrices  $U \in \mathbb{R}^{d,r}$  such that  $M \approx UV$ . To do so, one can try to minimise the function

$$L(U, V) = \frac{1}{d^2} \|M - UV\|^2$$

Implement this optimisation with:

1. Gradient descent
2. L-BFGS
3. Alternate minimisation strategy

First read in the image.

```

import imageio
from skimage.transform import resize

# load in file
M = imageio.imread("nus_logo.png").astype(float)
# transform to grayscale
M = onp.mean(M, axis = 2) / 255.
# resize
d = 100
M = resize(M, (d, d))

```

```
V_init = onp.random.normal(0, 1, size = (r, d)) /
    onp.sqrt(d*r)

plt.imshow(U_init @ V_init)
```




Figure 21: Initial matrix M

## → Gradient Descent

```
U = onp.copy(U_init)
V = onp.copy(V_init)

learning_rate = 10.
n_iter = 500
loss_history = []

for k in range(n_iter):
    gradient_u = grad_loss_u(U,V)
    gradient_v = grad_loss_v(U,V)

    # Update BOTH parameters in each iteration
    U = U - learning_rate * gradient_u
    V = V - learning_rate * gradient_v

    current_loss = loss(U,V)
    loss_history.append(current_loss)

plt.plot( loss_history )
plt.grid(True)
```




Figure 22: Loss plot (gradient descent)

```
plt.subplot(1,2,1)
plt.imshow(M, cmap="gray")

plt.subplot(1,2,2)
plt.imshow(U @ V, cmap="gray")
```




Figure 23: (Left) True image, (Right) Optimised image approx - gradient descent

## → Alternate Minimisation

Do alternate minimisation (i.e. minimise loss w.r.t.  $U$ , then w.r.t.  $V$ ) with L-BFGS. The problem is non-convex in terms of being a joint function of  $U$ ,  $V$ . Alternate minimisation is a faster alternative approach since the function is marginally convex.

First define functions that does a few steps of LBFGS with  $U$  and  $R$  fixed respectively. **Note:** LBFGS expects *vectors*,

not matrices so everything has to be flattened.

```
# Optimise V, done wrt U
def minimize_U_fixed(U, V, LBFGS_iter):
    """
    Given current U, V - fix U and optimise V with
    LBFGS
    Output optimised V (unflattened)
    """

    # 1) Function to unflatten V and output loss
    gradient
    def grad_v_wrapper(V_flat):
        V = V_flat.reshape(r, d)
        return onp.array(grad_loss_v(U, V)).ravel()

    # 2) Function to unflatten V and output loss
    def loss_wrapper(V_flat):
        V = V_flat.reshape(r, d)
        return loss(U, V)

    # 3) L-BFGS-B
    traj = scipy.optimize.minimize(
        fun = loss_wrapper, # function to minimise
        x0 = onp.array(V.ravel()), # initial guess
        method = 'L-BFGS-B',
        jac = grad_v_wrapper, # fn that computes gradient
        options = {"maxiter": LBFGS_iter} # max number of
        iterations
    )

    V_final = traj["x"].reshape(r, d)
    return V_final

# Optimise U, done wrt V
def minimize_V_fixed(U, V, LBFGS_iter):
    """
    Given current U, V - fix V and optimise U with
    LBFGS
    Output optimised U (unflattened)
    """

    # 1) Function to unflatten U and output loss
```

```

    gradient
def grad_u_wrapper(U_flat):
    U = U_flat.reshape(d, r)
    return onp.array(grad_loss_u(U, V)).ravel()

# 2) Function to unflatten U and output loss
def loss_wrapper(U_flat):
    U = U_flat.reshape(d, r)
    return loss(U, V)

# 3) L-BFGS-B
traj = scipy.optimize.minimize(
    fun = loss_wrapper, # function to minimise
    x0 = onp.array(U.ravel()), # initial guess
    method = 'L-BFGS-B',
    jac = grad_u_wrapper, # fn that computes gradient
    options = {"maxiter": LBFGS_iter} # max number of
        iterations
)

U_final = traj["x"].reshape(d, r)
return U_final

```

Start optimisation by looping through iterations - each iteration consists of optimising U and V (separately). They are optimised by further iterations of LBFGS.

```
U, V = onp.copy(U_init), onp.copy(V_init)
```

```
LBFGS_iter = 50
```

```
n_iter = 30
```

```
loss_history = []
```

```

for _ in range(n_iter):
    V = minimize_U_fixed(U, V, LBFGS_iter)
    U = minimize_V_fixed(U, V, LBFGS_iter)
    loss_history.append(loss(U, V))

```

```
plt.plot(loss_history, "-s")
```




Figure 24: Loss plot (alternate minimisation with LBFGS)

Plot out true image with approximated image.

```

plt.subplot(1,2,1)
plt.imshow(M, cmap="gray")

plt.subplot(1,2,2)
plt.imshow(U @ V, cmap="gray")

```




Figure 25: (Left) True image, (Right) Optimised image approx - alt min, LBFGS

## Function Reconstruction

In this type of optimisation problems, the goal is to recover the true y-values of the function. We are given data where the y-values modified.

Possible modifications:

1. y-values added with random noise
2. y-values go through another function (e.g. poisson)

### → Sine observations

Discussed in an example, likely in the realm of Tutorial 8 - regularised function reconstruction (given **only data points, approx function**). This is different from the previous example of "Function Approximation" since the parameters of a function are not found here. Instead, given  $\{x, z\}$  pairs, where  $z = y + \text{noise}$ , the goal is to **recover the y values**.

First, create samples drawn from the sine function with random noise added.

```

def func(x):
    """ one arbitrary function """
    return np.sin(x)

# Generate noise samples
n_samples = 100
sigma_noise = 0.3
x_list = onp.linspace(0, 10, n_samples)
z_true = func(x_list)
y_list = z_true + onp.random.normal(0,
    sigma_noise, n_samples)

plt.plot(x_list, y_list, "x")
plt.grid(True)

```




Figure 26: Scatterplot of data points (sine + noise)

---

```
reg = 5.
def loss(z):
    """
    loss function for regularised function
    reconstruction
    """
    difference = z[1:] - z[:(-1)]
    return np.mean( (z - y_list)**2 ) + reg *
        np.mean(difference**2)

grad_loss = jax.jit(jax.grad(loss))

# gradient descent
z = onp.copy(z_init)
learning_rate = 0.5
n_iter = 1000
loss_history = []

for k in range(n_iter):
    z = z - learning_rate*grad_loss(z)
    curr_loss = loss(z)
    loss_history.append(curr_loss)

plt.plot(loss_history)
plt.xlabel("Iteration")
```

---

```
plt.ylabel("Loss")
plt.grid(True)
```




Figure 27: Loss history (sine reconstruction)

---

```
plt.plot(x_list, y_list, "x")
plt.plot(x_list, z_true, "r-")
plt.plot(x_list, z, "b--")
plt.title("Function Reconstruction")
plt.grid(True)
```

---




Figure 28: Recovered sine values

### → Poisson observations

This is discussed likely for Tutorial 8. The true labels are now processed not by added noise but are inputs to a poisson distribution which outputs a random integer.

---

```
%matplotlib inline
from jax.config import config
config.update("jax_enable_x64", True) #important
for L-BFGS

import jax
import jax.numpy as np

import pylab as plt
import imageio
import os
import numpy as onp
from skimage.transform import rescale, resize,
    downscale_local_mean
import pandas as pd
```

---

First, generate the samples.

---

```

# list of coordinates
x_list = onp.linspace(0, 10, 100)

# arbitrary function
def func(x):
    return 10 + 7*np.sin(3*x)
z_true = func(x_list)

# list of poisson observations
y_list = []
for z in z_true:
    y_list.append(onp.random.poisson(z))
y_list = onp.array(y_list)

plt.plot(x_list, z_true, "r--")
plt.plot(x_list, y_list, "bo", alpha=0.5)
plt.grid(True)

```




Figure 29: Scatterplot of given values (blue dots) and actual function (red dotted line)

Define functions, negative loglikelihood. Note that output needs to be reparameterised (like in projected gradient descent), since poisson function can only take in positive values.

```
def neg_loglik(z):
```

```

"""
using it directly during optimisation may
through errors because z may become
negative
"""
regularization = np.mean( (z[1:] - z[:99])**2 )
log_lik = np.mean(-z + y_list * np.log(z))
return -log_lik + 0.1 * regularization

def neg_loglik_reparam(u):
    """
    u = log(z) <---> z = exp(u)
    """
    z = np.exp(u) # ensure input is positive
    return neg_loglik(z)

grad_neg_loglik_reparam =
    jax.jit(jax.grad(neg_loglik_reparam))

```

Optimise with basic gradient descent.

```

u = onp.zeros(100) # 100 data points given
n_iter = 1000
learning_rate = 0.1
loss_history = []
for k in range(n_iter):
    u -= learning_rate *
        grad_neg_loglik_reparam(u) # if just
        log-likelihood, need to maximise - so do
        the update with a plus (i.e. u = u + ...)
    loss_history.append(neg_loglik_reparam(u))

plt.plot(loss_history)
plt.grid(True)

```




Figure 30: Loss history (poisson)

To find the optimal set of values:

```

# compute z_MLE out of the u_MLE
z_MLE = np.exp(u)

plt.plot(x_list, z_true, "r--", label="Truth")
plt.plot(x_list, z_MLE, "g-", linewidth = 3,
        label="Reconstruction")
plt.plot(x_list, y_list, "bo", alpha=0.3,
        label="data")
plt.legend()
plt.grid(True)

```




Figure 31: Recovered function values

## Confidence Intervals

Discussed about the time of Tutorial 8, do so using Hessian at minimum (i.e. MLE).

Given some samples  $(y_1 \dots y_N)$ , fit a Gaussian  $N(\mu, \sigma^2)$  to it. We do not know the mean  $\mu$  and we do not know the standard deviation  $\sigma > 0$ .

Complete the following:

1. Write down the likelihood
2. Minimise the negative log-likelihood
3. Take care of the constraint  $\sigma > 0$  by re-parameterisation
4. Get the confidence intervals using the Hessian at the minimum

**Remark:** The density of a Gaussian distribution with mean  $\mu$  and standard deviation  $\sigma > 0$  is

$$\frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2\sigma^2}(x - \mu)^2\right)$$

First, generate some data.

---

```
mu_true = 1.
sigma_true = 2.
N = 100
y = onp.random.normal(mu_true, sigma_true, size =
    N) # NOTE: takes in s.d., not variance
H = plt.hist(y, bins = 20)
```

---




Figure 32: Histogram of true data (N = 100)

**Question 1. Write down the likelihood.**

**Question 2. Minimise the negative log-likelihood.**

**Question 3. Take care of the constraint  $\sigma > 0$  by re-parameterisation.**

---

# IMPORTANT: know how to derive this

```
def mean_negloglik(mu, sigma):
    return np.mean(np.log(sigma) +
        (y-mu)**2/(2*sigma**2))
```

---

In order to deal with the constraint  $\sigma > 0$ . Introduce instead  $u = \log \sigma$ . In other words, we are looking for the parameter  $x = (\mu, u) \in \mathbf{R}^2$ .

---

# Define functions, gradient

```
def loss(x):
    """ loss == mean negative log likelihood """
    mu = x[0]
    u = x[1]
    sigma = np.exp(u)
    return mean_negloglik(mu, sigma)
```

```
grad_loss = jax.jit(jax.grad(loss))
```

# NECESSARY FOR LBFGS

```
def gradient_wrapper(x):
    """ computes the gradient of function at x,
        makes sure that output is a numpy array"""
    return onp.array(grad_loss(x))

loss_history = [] # save the loss trajectory

def save_traj(x):
    """ function that saves a few statistics for
        later analysis"""
    loss_history.append(loss(x))

# Initialise values
x_init = onp.array([0.,0.])
grad_loss(x_init) # test: [-1.29710711, -5.3042877 ]

# Optimise
traj = scipy.optimize.minimize(
    fun = loss, # function to minimise
    x0 = onp.array(x_init), # initial guess
    method = 'L-BFGS-B',
    callback = save_traj, # function used to save some
        results for later plotting
    tol = 10**-10, # NOTE: important to set tolerance,
        esp for images
    options={"maxiter":100} # maximum number of
        iterations
)

# Plot loss
plt.plot(loss_history, "-s")
plt.title("mean negative log likelihood")
plt.grid(True)

# Optimal values
mu_MLE, u_MLE = traj["x"]
sigma_MLE = onp.exp(u_MLE)
mu_MLE, sigma_MLE # (1.2971070934835425,
    2.1498373989917465)
```

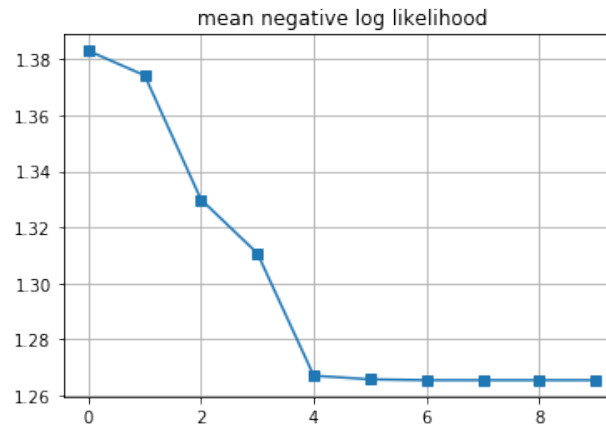


Figure 33: Loss plot

#### Question 4. Get the (90%) confidence intervals using the Hessian at the minimum.

The confidence interval for the estimated parameters can be estimated by computing the Hessian  $\mathbf{H}$  at the *MLE*. The **inverse** Hessian is an approximation of the covariance structure of the parameter estimates.

**Remark:** Get 90% C.I. using  $CI = estimate \pm 2 * s.d..$

The approximate covariance structure is the inverse of the hessian of the negative log-likelihood (not the mean log-likelihood), hence the use of sum over the individual likelihoods for each observation. A good/intuitive way to understand is that the more data you have, the lower the confidence interval should be. Using the average log-likelihood will not shrink the confidence interval as more data is obtained.

```
def true_negloglik(z):
    mu = z[0]
    sigma = z[1]
    return np.sum(np.log(sigma) +
        (y-mu)**2/(2*sigma**2))

# Compute the Hessian
hess_neg_loglik = jax.hessian(true_negloglik)
```

```
# -----

# Compute the Hessian at the MLE
param_MLE = onp.array([mu_MLE, sigma_MLE])
hessian = hess_neg_loglik(param_MLE) # a 2x2 array

# Compute the inverse of Hessian to approximate
    the covariance
covariance_MLE = onp.linalg.inv(hessian) # a 2x2
    array

# -----

# Compute marginal standard deviation
# NOTE: standard deviation of parameters are roots
    of diagonal elements (i.e. variances)
std_mu = onp.sqrt(covariance_MLE[0, 0])
std_sigma = onp.sqrt(covariance_MLE[1, 1])

print("MLE estimates: \t MLE(mu)={0:2.2f} \t
    MLE(sigma)={1:2.2f}".format(std_mu, std_sigma))
print("90% CI for mu is about
    [{0:2.2f},{1:2.2f}].format(mu_MLE - 2*std_mu,
        mu_MLE + 2*std_mu))
print("90% CI for sigma is about
    [{0:2.2f},{1:2.2f}].format(sigma_MLE -
        2*std_sigma, sigma_MLE + 2*std_sigma))

MLE estimates:  MLE(mu)=1.18  MLE(sigma)=2.15
90% CI for mu is about [0.87,1.73]
90% CI for sigma is about [1.85,2.45]
```

Plot the joint distribution of the Gaussian approximation of the parameters at the *MLE*. i.e. Using the MLE estimates, sample data from Gaussian distribution. Indicate the true and MLE pair of  $(\mu, \sigma)$  using coordinates.

```
samples_gaussian =
    onp.random.multivariate_normal(param_MLE,
        covariance_MLE, size=10000)
```

```
plt.scatter(samples_gaussian[:,0],
            samples_gaussian[:,1], alpha=0.2, marker="+")
plt.scatter(mu_MLE, sigma_MLE, c="red",
            marker="^", s=100, label="MLE")
plt.scatter(mu_true, sigma_true, c="purple",
            marker="s", s=100, label="True")
plt.legend()
plt.grid(True)
plt.xlabel(r"$\mu$")
plt.ylabel(r"$\sigma$")
```

---




Figure 34: MLE estimate of  $\mu$  and  $\sigma$  compared to true values