

Docker

(Learning from KodeKloud's YouTube video *Docker for Beginners: Full Course*)

Why do we need Docker?

→ Issues

1) Compatibility/Dependency

- Using many different components in an end-to-end application
- Version of operating system, dependent libraries etc in architecture
- Compatibility matrix issue

2) Long setup time

- For new developers
- Many instructions and command to set up the application

3) Different developer/test/production environments

- Due to difference preferences in OS, versions etc
- Need to ensure that the application is running the same way in different environments

→ Features

- Runs each component in a separate container with its own dependencies, libraries etc
- Manages these components
- Same VM or OS
- Build Docker configuration once to set up the environment

→ Benefits

- Solves compatibility issue
- Allows the modification of components without affecting the other components
- Lightweight, boots up faster
- Less isolations, more resources are shared between the containers

Containers

- Completely isolated environments (e.g. file system etc)
- Have their own processors, network interfaces, mounts etc
- All share the same OS kernel (Linux, Windows etc)
- Container OS needs to match host kernel

→ How they work

- Not mean to host an operating system
- Meant to run a specific task or process
- For example, host an instance of a web server, application server, database or compute/analysis something
- Once task completed, container exits (only lives as long as processes in it are running)
 - If service in container is stopped/crash, container exits
- Does *not* listen to a standard input (i.e. input from console), unlike a regular CLI app
- Runs in a non-interactive mode

→ How to users access the application

- By using the port which the application is listening on
- To access from web browser:
 - 1) Use internal IP of Docker container (only accessible within Docker host)
 - 2) Do *port mapping*

→ Running a container for an OS

More precisely if image is just a base image for other images (e.g. Ubuntu),

- Runs an instance of image
- Exits immediately
- Container status will be shown as "Exited"

Why?

- By default, Docker does not attach a terminal to a container when it is run
- Bash does not find a terminal and exits
- Process started when terminal created (i.e. Bash) finishes, therefore container exits

Operating systems

- Consists of 2 things: OS kernel + set of software

OS Kernel	Set of software
- Responsible for interactions with underlying hardware	- Makes operating systems different - User interface, drivers, compilers, file managers, developer tools etc

Docker Image		Docker Container	
<ul style="list-style-type: none">- Package/template/plan- Used to create one or more containers- Can be created with a Docker file of requirements from both Developers and Operations (based on the Developer Guide and instructions for the application)- Can be deployed in production- Examples: mongoDB, nodejs, redis etc		<ul style="list-style-type: none">- Running instance of image- Isolated, have their own environments, set of processors	
Volume mapping			
<p>→ How data is persisted in a Docker container</p> <ul style="list-style-type: none">- Example: Running a MySQL container<ul style="list-style-type: none">- Creates databases and tables- Data stored in /var/lib/mysql in container- Docker container has its own isolated filesystem (any changes in files happen in container)- Directory name in “/var/lib/docker/” folder is the same as container ID for files related to container		docker run mysql	
<ul style="list-style-type: none">- All data in container is gone		docker stop mysql docker rm mysql	
<ul style="list-style-type: none">- To keep data, map directory outside the container on Docker host to directory inside container- ** Implicitly mounts external directory to a folder inside the docker container- All data will be stored at external volume in the outside Docker directory- Remains even if container is deleted		docker run -v <docker host dir>:<container dir> <container name>	
Creating an image			
1) Create Dockerfile and write down instructions for setting up your application in terms of CLI commands you would have used (like a bash script)			
Example	FROM Ubuntu RUN apt-get update RUN apt-get install python RUN pip install flask RUN pip install flask-mysql COPY . /opt/source-code ENTRYPOINT FLASK_APP=/opt/source-code/app.py flask run		
2) Build image using docker build command and tag the application version Layer 1. Base Ubuntu layer Layer 2. Changes in apt packages Layer 3. Changes in pip packages Layer 4. Source code Layer 5. Update Entrypoint of image with “flask” command		<ul style="list-style-type: none">- Docker builds images in an <i>layered architecture</i>- Each line of instruction creates new layer in image- <i>Docker build output:</i> Can see each step and output of each command/layer/task	
	docker build Dockerfile -t mmumshad/my-custom-app		
3) Make local image available on Docker registry <ul style="list-style-type: none">- Meant to be small and lightweight			
	docker push mmumshad/my-custom-app		
If a layer fails during the image build process			

<ul style="list-style-type: none">- After fixing issue and running “docker build”, reuse previous layers from cache and continue to build the remaining layers- Faster rebuilding of image, no need to start from first layer of build process- Helpful if updating source code of application (only layers above updated layers need to be rebuilt)			
DockerFile			
Argument			
CMD	Overrides default command (specified in image’s “CMD”) to run when container starts		
ENTRYPOINT	Appends param to command indicated in “ENTRYPOINT”		
Arguments CMD and ENTRYPOINT			
- An ENTRYPOINT command can be a bash script (.sh)			
Examples of Dockerfile	<pre># Overrides default command FROM Ubuntu # Shell format CMD sleep 5 # or JSON format: CMD ["command", "param1"] CMD["sleep", "5"] ## CMD: command line params <u>override</u> completely</pre>	<pre>FROM Ubuntu # Specify program to run when container starts ENTRYPOINT["sleep"] ## ENTRYPOINT: command line params get <u>appended</u> # > docker run <u>ubuntu-sleeper</u> 10 # is the same as # > docker run <u>ubuntu</u> sleep 10</pre>	<pre># Always specify in <u>JSON</u> format FROM Ubuntu ENTRYPOINT["sleep"] # <u>Default value</u> if no value is specified in command line, if have then that value overrides this CMD["5"] # sleep 5 sec # > docker run ubuntu-sleeper # sleep 10 sec # > docker run ubuntu-sleep 10</pre>
Examples in CLI	<pre># Modify ENTRYPOINT command <u>during runtime</u> # sleep2.0 is the new command docker run --<u>entrypoint</u> ubuntu-sleeper 10</pre>		
Networks			
<p>→ Default networks (when Docker is installed)</p> <ul style="list-style-type: none">1) Bridge2) Null/None3) Host <p>- Docker has a built-in DNS server</p> <ul style="list-style-type: none">- Helps containers resolve each other using the container names- Always runs at address 127.0.0.11 <p>→ How are containers isolated within host?</p> <ul style="list-style-type: none">- Use network name spaces- Create separate network name space for each each container- Use virtual ethernet pairs to connect containers together			
Bridge	<ul style="list-style-type: none">- Default network that a container gets attached to- Private internal network- Created by Docker on the host- Containers attached to this network usually get IP address in the range 172.17 series (172.17.0.1 ... 172.17.0.5 etc)		

	<ul style="list-style-type: none"> - Containers can access each other using internal IP if required - To access these containers to outside world, Method 1) Map ports on containers to ports on docker host - By default, only one internal network is created 	
None	Containers: <ul style="list-style-type: none"> - Are not attached to any network - Does not have access to external network or other containers - Run in an isolated network 	
Host	<ul style="list-style-type: none"> - To access these containers to outside world, Method 2) Associate container to host network - Take out any network isolation between Docker host and Docker container (no need port mapping, container uses the host's network) - No longer able to run multiple containers on same host and same port (ports now common to <i>all</i> containers in host network) 	
Storage drivers		
<ul style="list-style-type: none"> - Responsible for <ul style="list-style-type: none"> - Maintaining layered architecture - Creating a writable layer - Moving files across to enable copy-on-write etc - Selection of driver depends on underlying OS being used <ul style="list-style-type: none"> - Ubuntu: AUFS driver - Different driver, different performance and stability characteristics - Choose base on needs - Common drivers <ul style="list-style-type: none"> - AUFS, ZFS, BTRFS, Device Mapper, Overlay, Overlay2 		
Docker compose		
<ul style="list-style-type: none"> - When setting up a complex application such as running (<i>docker run ...</i>) multiple services - Create a configuration file in <i>yaml</i> format, put together different services and options specific to the file - Only applicable to running all containers on a <i>single</i> docker host 		docker-compose.yml
→ How to do so <ul style="list-style-type: none"> - Individually <i>docker run</i> the necessary containers <ul style="list-style-type: none"> - Naming containers is important - Link all the containers together (tell them to use each other specifically) - Use command-line option to link container to container it is depending on 		# Bring up entire app stack
Structure of <i>yaml</i> file: <ul style="list-style-type: none"> - Dictionary of container names - Under each item, specify image to use - ** Basically translating all the docker run commands and their parameters into this file 		docker-compose up
Docker registry		
<ul style="list-style-type: none"> - Central repo of Docker images - For example, image <code>nginx</code> <p style="text-align: center;">image: nginx</p> <p>actually refers to</p> <p style="text-align: center;">image: nginx/nginx</p> <p style="text-align: center;">(user or account/image or repository)</p> <p>and is stored at Docker hub</p> <p style="text-align: center;">image: docker.io/nginx/nginx</p> <p style="text-align: center;">(registry/user/image)</p>		
Private registries <ul style="list-style-type: none"> - Good for internal private in-house applications - Accessible only with credentials - If not logged in, Docker will say that image cannot be found 		docker login private-registry.io

	docker run private-registry.io/apps/internal-app
Docker engine	
<ul style="list-style-type: none"> - Docker engine = host + docker installed on it - Related to docker architecture <ul style="list-style-type: none"> - how it runs applications in isolation - how it runs under the hood - What is installed when Docker is installed <p>1) Docker CLI</p> <ul style="list-style-type: none"> - Command-line interface - Uses REST API to interact with Docker Daemon - Does not need to be on the same host, can be on another system like a laptop <p>2) REST API</p> <ul style="list-style-type: none"> - API interface that programs and used to talk to daemon and provide instructions - Own tools can be created using this API <p>3) Docker Daemon</p> <ul style="list-style-type: none"> - Background process that manages Docker objects (e.g. images, containers, volume and networks etc) 	
Containerisation	
<ul style="list-style-type: none"> - Docker uses <i>namespaces</i> to isolate workspaces - Workspaces such as <ul style="list-style-type: none"> - Processing IDs - Network - Inter-process communication - Mounts - Unix time-sharing systems - Each workspace created in their own namespace, provides isolation between containers 	
<p>Process ID (PID) namespaces</p> <ul style="list-style-type: none"> - When Linux system boots up, starts with one process with a process ID of one <ul style="list-style-type: none"> - Root process - Kicks start all the other processes - By the time system starts up completely, there are a handful of processors running - Check processes with <code>ps</code> command - Processors are unique (no two processors have the same process ID) - A container/child system within Linux system <ul style="list-style-type: none"> - Thinks it is an independent system of its own - Has its own set of processes - No hard isolation between containers and underlying host, so processors running within container are actually running on host - Unique PID - With PID namespaces, each process can have multiple process IDs, associated with it - Process that starts in container gets next available PID on underlying system 	<p><u>Linux system</u></p> <p>PID: 1 PID: 2 PID: 3 PID: 4 (same)</p> <p><u>Child system (container)</u></p> <p>PID: 1 PID: 2 (same)</p>
Container Orchestration	
<p>→ Situation</p> <ul style="list-style-type: none"> - If number of users increase, and need to manually run new instances of container multiple times - Need to deploy additional instances and keep watch of load and performance of application - Need to also keep watch if a container fails, or if the host crashes <p>→ Solution</p> <ul style="list-style-type: none"> - Avoid manual monitoring to some extent - A set of tools and scripts that can help host containers in a production environment - Typical solution: multiple Docker hosts that can host containers <ul style="list-style-type: none"> - Even if one fails, application is still accessible 	

<ul style="list-style-type: none">- Deploy 100s, 1000s of instances of applications with a single command- Some solutions:<ul style="list-style-type: none">- Scale up when demand increase, scale down when demand decrease (scaling)- Automatically add additional hosts to support user load (clustering)- Advanced networking between containers across different hosts- Load balancing across different hosts- Provide sharing storage between host, support management and security within cluster <p>→ Well-known solutions today</p> <ul style="list-style-type: none">- Docker Swarm<ul style="list-style-type: none">- lacks advanced complex scaling features- Kubernetes<ul style="list-style-type: none">- a bit difficult to set-up- provides a lot of options to customize deployments- has support for many different vendors (e.g. public cloud service providers)- popular- MESOS ("mez-o-s")<ul style="list-style-type: none">- difficult to set-up and get started, supports many advanced features	

Docker Swarm (docker swarm)

- Combine multiple Docker machines together into a single cluster
- Will take care of distributing services/application instances into separate hosts for high availability and balancing across different systems and hardware
- Requires multiple hosts with docker installed on them
- Designate one to be *swarm manager/master* and the rest as *workers/slaves/nodes*

on docker swarm manager host

docker swarm **init** ... # initialize manager

on workers

docker swarm **join** --token <token> # command provided by manager during init

Docker service

- Key component of swarm orchestration
- Definition: one or more instances of a single application/service that runs across the nodes in the swarm cluster

run on the manager node**

indicate number of instances of image to run across the cluster

docker **service create** --replicas=<int> <image name>

can use arguments like --network, -p for ports etc

Kubernetes (kubectl)

- Instances can automatically scale up or down based on user load
- Can create more instances than Docker Swarm
- kubernetes open architecture provides support for
 - many storage networks
 - a variety of authentication and authorization mechanisms
- All major cloud service providers have native support for Kubernetes
- **kubernetes uses Docker hosts to host applications in the form of Docker containers**
- or alternative containers (e.g. rkt, cri-o)

→ kubernetes architecture

- kubernetes cluster consists of a set of *nodes*
 - Even if one node fails, application still accessible from other nodes
- A *node*
 - physical or virtual machine (a *worker machine*)
 - where Kubernetes software/set of tools are installed
 - where containers are launched by kubernetes
 - if node where application is running on fails, node goes down

→ Managing cluster

- Store information about cluster
- How are nodes monitored
- When node fails, how to move workload of failed node to another worker node
- Done by a *master node*
- Has Kubernetes control plane components installed
- Watches over nodes in cluster
- Responsible for actual orchestration of containers on the worker nodes

→ Installing kubernetes

- Multiple components installed

1) API server

- For front-end for user management devices, CLI interfaces to talk to API, interact with cluster

<p>2) etcd server</p> <ul style="list-style-type: none"> - Reliable key-value store - Stores all data to manage the cluster - Stores all information about multiple nodes in a cluster in a distributed manner (considering that there are multiple clusters as well) - Responsible for implementing a law within the cluster to ensure no conflicts between master nodes <p>3) kubelet service</p> <ul style="list-style-type: none"> - Agent that runs on each node in the cluster - Responsible for making sure that the containers are running on the nodes as expected <p>4) Container Runtime</p> <ul style="list-style-type: none"> - Underlying software that is used to run containers - E.g. an engine like Docker <p>5) Controllers</p> <ul style="list-style-type: none"> - The “brain” behind orchestration - Responsible for noticing and responding when nodes, containers or endpoints goes down - Makes decisions to bring up new containers in cases of crashed nodes <p>6) Scheduler</p> <ul style="list-style-type: none"> - Responsible for distributing work/containers across multiple nodes - Looks for newly created containers and assigns them to nodes 	
<p>→ kubernetes CLI</p> <ul style="list-style-type: none"> - Or called “cube control tool” - Used to deploy or manage applications on cluster to get cluster-related information, status of nodes in cluster etc 	
<p># Examples</p> <p># deploy application on cluster</p> <pre>kubectl run hello-minikube</pre> <p># view information about cluster</p> <pre>kubectl cluster-info</pre> <p># list all nodes part of cluster</p> <pre>kubectl get nodes</pre>	
<p># Examples</p> <p># run many instances of the same application</p> <pre>kubectl run --replicas=1000 my-web-server</pre> <p># scale up the number of instances (based on user load)</p> <pre>kubectl scale --replicas=2000 my-web-server</pre> <p># upgrade 2000 instances in a rolling upgrade fashion one-at-a-time</p> <pre>kubectl rolling-update my-web-server --image=web-server:2</pre> <p># roll back images if something goes wrong</p> <pre>kubectl rolling-update my-web-server --rollback</pre> <p># test new features of application, upgrade only some instances (A/B testing)</p> <p>...</p>	

Term	Meaning	Notes
Docker hub/store	public Docker repository	
Load-balancer		
DevOps	Developer + Operations team	
Hypervisor	<ul style="list-style-type: none"> - Also known as “virtual machine monitor” (VMM) - Computer software, firmware or hardware that creates and runs virtual machines 	
Host machine	A computer on which a <i>hypervisor</i> or more runs	
Guest machine	Each virtual machine running on the host machine	
Attached mode	<ul style="list-style-type: none"> - Attached to console/standard output of Docker container - Able to see output of application on console - Won't respond to user inputs - Stop container and exit: CTRL+C 	When running a container with docker run
Detached mode	- Container runs in back end	
Supported tags	<ul style="list-style-type: none"> - Associated with an image at dockerhub.com - Each version of a software can have multiple short and long tags 	
IP of Docker host	<ul style="list-style-type: none"> - Local host port/IP - “<i>application is available on port ... on the host</i>” 	
IP of Docker container	<ul style="list-style-type: none"> - Internal IP - Assigned by default (each container has one) - Only accessible within Docker host - “<i>application listens on port ...</i>” 	http://...
Port mapping/publishing	<ul style="list-style-type: none"> - Routes all traffic on local port to internal port - ** Note: Cannot map to same port on Docker host more than once <p>Benefits</p> <ul style="list-style-type: none"> - Run multiple instances of application and map them to different ports on docker host - Can run as many applications and map them to as many ports desirable 	docker run -p <local port>:<internal port> <image name>
Volume mapping	<ul style="list-style-type: none"> - Saves all data in container to external directory on Docker host - Refer to notes above 	
Logs	Contents written to standard output of the container	
Dockerfile	<ul style="list-style-type: none"> - Text file written in a specific format that Docker can understand - In an <i>instruction-and-argument</i> format - All Dockerfiles must start with a “<i>from</i>” instruction - Possible naming convention if there are multiple Dockerfiles, “Dockerfile-<app name>” <p>Instruction: FROM, RUN, COPY, ENTRYPOINT, EXPOSE, WORKDIR, CMD etc</p> <p>Argument: <bash command></p>	FROM <base OS/image> RUN <command to run on base image> # Copy files from local system to Docker image COPY <src code dir> <dest src code dir in Docker> # Port which application is run within container EXPOSE <port no.> # Specify a command to run when image is a container (command to run app) ENTRYPOINT ...
Docker build output	<ul style="list-style-type: none"> - When running “docker build” command - See various steps involved and their task output 	

[illegible]

Commands	Purpose	Notes
docker version	List version of Docker used by host	
docker -H=<remote docker engine address>:<port>	Run a container using a remote docker host	
docker service create -- replicas=<int> <image name>	Create multiple docker hosts for containers	- Helps prevent inaccessible application if a Docker host crashes
Example Command		Explanation
docker -H=remote-docker-engine:2375		
docker service create --replicas=100 nodejs		
docker images		
docker pull <app name>	- Only pull/download image and store on host - Does <i>not</i> create a container	
docker images	List all available images on host	Output column names: REPOSITORY (image name) TAG IMAGE ID CREATED SIZE
docker rmi <image name>	- Remove images for good - Make sure no containers are running off of that image first (<i>delete all dependent containers</i>) - Image name is in the form "REPOSITORY:TAG"	- Remove multiple images by adding more image names
docker build Dockerfile -t <tag name for image>	- Build image from application ("dockerise app") - Creates image locally on system	
docker push <image name>	Make local image available on Dockerhub registry	Example of image name: account_name/app_name
docker history <image name>	See layers in image	
Example		Explanation
docker build Dockerfile -t mmumshad/my-custom-app		Build an image from an application with a tag of the app version
docker build -t webapp-color .		Build a docker image using the Dockerfile (currently in directory of file) with no tag specified - Fullstop (.) represents Dockerfile
docker build -t webapp-color:lite .		Build an image, name it <i>webapp-color</i> and tag it <i>lite</i> (use colon)
mmumshad/my-custom-app		An example of an image name
docker run		
docker run <image name>	- Run a container from an image - If image not present locally, it will be pulled from Docker hub and stored - Runs in attached mode (foreground) More precisely if image is just a base image for other images (e.g. Ubuntu), - Runs an instance of image - Exits immediately - Container status will be shown as "Exited"	
docker run -d ...	- Run a container from an image in detached mode (background)	

docker run -i ...	<ul style="list-style-type: none"> - To enable application to read inputs (interactive mode) - Map standard input of host to Docker container - Short for “interactive” 	Note: Application prompt on terminal will be missing (e.g. “Please enter your name.”)
docker run -t ...	<ul style="list-style-type: none"> - Enable output by application as prompts - Attach terminal to container - Short for “pseudo terminal” 	
Mapping		
docker run -p <local host port>:<internal port> ...	<ul style="list-style-type: none"> - ** Map local host port to internal port (i.e. map port <i>inside</i> Docker container to a free port on the Docker host/expose application on hosts port) - All traffic on local port will be routed to internal port 	- Allows access for users outside Docker host
docker run -v <docker host dir>:<container dir> <container name/ID>	<ul style="list-style-type: none"> - ** Volume mapping - Store data of application, prevent its deletion upon stopping of container 	
Storage		
docker run --mount type=<volume/bind>, source=<file path>, target=<file path> <image name>	<ul style="list-style-type: none"> - ** Volume/bind mounting - Allows data to persist by storing data written by container into a read-write folder 	
docker compose		
docker run --link <container depended upon>:<host name> <image name>	<p>Allows a container to link to a container it needs/depends on</p> <ul style="list-style-type: none"> - Deprecated feature and may be removed in the future in Docker 	
Networking		
docker run <image name> --network=<network name>	Associate container with non-default network	<p>Network names are</p> <ul style="list-style-type: none"> - “Bridge” (default) - “none” - “host”
docker Engine		
docker -H=<engine address>:<port> run ...	Run a container on a remote engine host	
docker run --cpus=<0-1> ...	Restrict CPU access to container	
docker --memory=<limit> ...	Restrict memory usage by a container	
docker run -e <var name>=<value>	<ul style="list-style-type: none"> - Set/configure value to environment variable - Value will be passed into application and can be used 	
docker run --name <given name> <image name>	Run a container from an image and name it	
docker run <image name> :<version no.>	<ul style="list-style-type: none"> - Run another version of an image (e.g. older ver, another tagged version) - Specify version to use in a tag (identified with a colon) - (Default) No tag means latest tag (i.e. version) 	
docker run <image name> <command>	<ul style="list-style-type: none"> - Override default command in “CMD” instruction of image (command to run when container starts) 	docker run ubuntu sleep 5

<pre>docker run --entrypoint <new command> <image name> [<params...>]</pre>	<ul style="list-style-type: none"> - Specify a different command to start the container - When container starts, it runs the command 	
	Override command specified in "ENTRYPOINT" during runtime	
Example Command		Explanation
docker run ubuntu sleep 5		Container sleeps for 5 seconds
docker run -it <image name>		In interactive mode and attached terminal to container
docker run -p 80:5000 <image name>		Routes traffic from local port 80 to internal port 5000 (users access with <IP>:<local port> like http://192.168.1.5:80)
docker run -v /op/datadir:/var/lib/mysql mysql		Map data directory in docker host to directory (/var/lib/mysql) in container
docker run -e APP_COLOR=blue simple-webapp		Set environment variable APP_COLOR as blue for application to have blue background
docker run -e MYSQL_ROOT_PASSWORD=db_pass123 mysql		
docker run python:3.6 cat /etc/*release*		Find out base OS used by image
docker run Ubuntu --network=none		Associate container with another network other than "Bridge"
docker -H=10.123.2.1:2375 run nginx		Run nginx container on a remote docker host
docker run --cpus=.5 ubuntu		Ensure that container does not take up more than 50% of host CPU at any given time
docker run --memory=100m ubuntu		Limit the amount of memory container can use to 100 MB
docker containers		
docker attach <container ID/name>	Attach running detached container	For container ID, providing the first few characters is enough
docker start <container ID/name>	Start a stopped container	
docker ps	<ul style="list-style-type: none"> - List all running containers and their basic information - Container ID and name (randomly created by Docker) - Name of image - Current status (<i>Up</i>, <i>Exited</i>) 	Output column names: CONTAINER ID IMAGE (name of image) COMMAND CREATED STATUS PORTS (will show ports published on the container) NAMES (name of container)
docker ps -a	- List all running, previously stopped and exited containers (all present containers)	
docker inspect <container ID/name>	<ul style="list-style-type: none"> - Gives all details about a specific container - E.g. states, mounts, configuration, data, network settings etc - Returns all details in a JSON format - Inspect values of environment variables under "Config": "Env" section (of a running container) 	
docker stop <container ID/name>	Stop a container	- Stop multiple containers by adding more ID/names
docker rm <container ID/name>	<ul style="list-style-type: none"> - Remove a stopped/exited container for good - Success: prints container ID/name 	- Remove multiple containers by adding more ID/names

docker exec <running container ID/name> <command>	Execute command on running container	
docker logs <container ID/name>	See logs of container running/which ran in <i>detached mode</i> (background)	
Example Command		Explanation
docker exec mysql-db mysql -pdb_pass123 -e 'use foo; select * from myTable'		View information in database where container is mysql-db with password db_pass123
networking		
docker run <image name> -- network=<network name>	Associate container with non-default network	Note: duplicated from <i>docker run</i> section
docker run --link <src container name>:<container alias name>	<ul style="list-style-type: none"> - Allow container being created to be linked to another container ("<i>src container</i>") - Source container will be associated with alias name and can be referenced through it 	Note: Linking is a LEGACY feature - Use the above command, connect containers to the same network to link them
docker network create \ --driver <network> \ --subnet <IP> --gateway <IP> <custom isolated network name>	<ul style="list-style-type: none"> - Create a new network - Helps to isolate internal network (i.e. only certain containers share certain networks) - Not all fields compulsory 	docker network create \ --driver bridge \ --subnet 182.18.0.0/16 custom-isolated-network
docker network ls	List all networks	
docker network inspect <network>		
docker inspect <container ID/name>	<ul style="list-style-type: none"> - See network settings and IP addressed assigned to existing network container - In the JSON, inspect "Networks": "<bridge/none/host>" and "IPAddress", "MacAddress" 	Not guaranteed that the container will get the same IP address when system reboots
Connect containers of the same internal network	Use container name instead of IP address	
Example Command		Explanation
mysql.connect(172.17.0.3)		NOT recommended as IP address may change
mysql.connect(mysql)		Recommended , when web container is trying to connect to database container name " <i>mysql</i> " with internal IP address 172.17.0.3
docker network inspect bridge		Gives details and information such as the subnet configured on the <i>bridge</i> network
storage		
<ul style="list-style-type: none"> - Mounting direction is <folder mounting>:<folder mounted on> - Preserve data created by container to image (e.g. database) - Once mounted, any changes will be stored and just do <i>volume mapping</i> to get back the data when restarting a new instance of the container with <i>docker run -v ...</i> (e.g. if current container, database crashed) 		
Method 1 (Step 1)		
docker volume create <folder name>	Preserve data created by container to image/ add persistent volume to container	<ul style="list-style-type: none"> - New folder created /var/lib/docker aufs containers image volumes <folder_name>

Method 1 (Step 2)		
docker run -v <folder name>: /var/lib/<image name> <image name>	Mount folder created by <i>docker volume create</i> onto folder in image directory - Executed after <i>docker volume create</i> - Called volume mounting : mounts volume from /volumes directory	- If this command is not run before the <i>docker run</i> command, docker will automatically create the folder
Method 2		
docker run -v <complete file path>:<image dir> <image name>	- Alternative : not store data on default /var/lib/ folder - Called bind mounting : mounts a directory from any location onto docker host	- Provide complete file path of folder to mount onto container
Method 2 (alternative)		
docker run --mount type=<volume/bind>, source=<file path>, target=<file path> <image name>	- Alternative : better recent alternative to “-v” command	
Example Command		Explanation
docker volume create data_volume		- Run codes in succession
docker run -v data_volume:/var/lib/mysql mysql		
docker run -v /data/mysql:/var/lib/mysql mysql		- Method 2
docker run --mount type=bind, source=/data/mysql, target=/var/lib/mysql mysql		- Better alternative of Method 2
docker compose		
- Levels: 1) container names 2) properties (e.g. image, ports, links, depends_on, networks, environment etc) 3) values to properties		
docker run --link <container depended upon>:<host name> <image name>	Allows a container to link to a container it needs/depends on - Deprecated feature and may be removed in the future in Docker	Example of Python application that requires a host: def get_redis(): if not hasattr(g, 'redis'): g.redis = Redis(host="redis", db=0, socket_timeout=5) return g.redis
docker-compose up	Brings up entire application stack, based on docker-compose.yml - Aggregates the output of each container	- When command exits, all containers are stopped
Example Command		Explanation
docker run -d --name=vote -p 5000:80 --link redis:redis voting-app		Creates an entry in /etc/hosts file in voting-app container, adds an entry with host named redis, with internal IP of redis container
docker run -d --name=worker --link db:db --link redis:redis worker		In worker application: try { Jedis redis = connectToRedis("redis"); Connection dbConn = connectToDB("db"); System.err.println("Watching vote queue"); ...

docker-compose up -d	Starts the containers in the background and leaves them running
docker compose format (VER 1)	
- Disadvantages: - Unable to deploy application on networks other than default <i>bridge</i> network - Cannot specify dependency/start-up order (e.g. database needs to be set up first for other app to run properly)	
docker-compose.yml <pre> redis: image: redis db: image: postgres:9.4 vote: image: voting-app ports: - 5000:80 links: - redis result: image: result-app ports: - 5001:80 links: - db worker: image: worker links: - redis - db </pre>	Rather than running one by one: <pre> docker run -d --name=redis redis docker run -d --name=db postgres:9.4 docker run -d --name=vote -p 5000:80 --link redis:redis voting-app docker run -d --name=result -p 5001:80 --link db:db result-app docker run -d --name=worker --link db:db --link redis:redis workers </pre> <p>- Note that here, db:db == db</p>
OR ... build: ./vote ...	** Instruct Docker to do a build instead of pulling an image - When image is your own application - Use file path to application source code and <i>Dockerfile</i>
docker compose format (VER 2)	
- No longer display stack information as directly as before (all encapsulated in <i>services</i> section) - Can define networks to use	
docker-compose.yml <pre> version: '2' services: redis: image: redis networks: - back-end db: image: postgres:9.4 networks: - back-end vote: image: voting-app ports: - 5000:80 </pre>	Shift all of Version 1 under <i>services</i> 1) ** Specify version of docker compose file if using version 2 or higher - Write version number as string 2.1) No need to use link in version 2 and higher since docker will automatically create a <i>dedicated</i> bridge network for this application - All containers are attached to this network 2.2) Dedicate networks to use - Front-end network: traffic for users - Back-end network: traffic for applications 3) Add dependencies with depends_on property - Indicate which container is dependent on

links:

- redis

depends_on:

- redis

networks:

- front-end

- back-end

result:

image: result-app

ports:

- 5001:80

links:

- db

networks:

- front-end

- back-end

worker:

image: worker

links:

- redis

- db

networks:

- back-end

networks:

front-end:

back-end:

docker compose format (VER 3)

- Similar to Version 2
- Some options removed or added
- Comes with support for *docker swarm*

docker-compose.yml

version: '3'

services:

redis:

image: redis

db:

image: postgres:9.4

vote:

image: voting-app

ports:

- 5000:80

links:

- redis

depends_on:

- redis

result:

image: result-app

ports:

